



Image Analysis and Object Recognition

Exercise 5

SIFT Scale Invariant Feature Transform

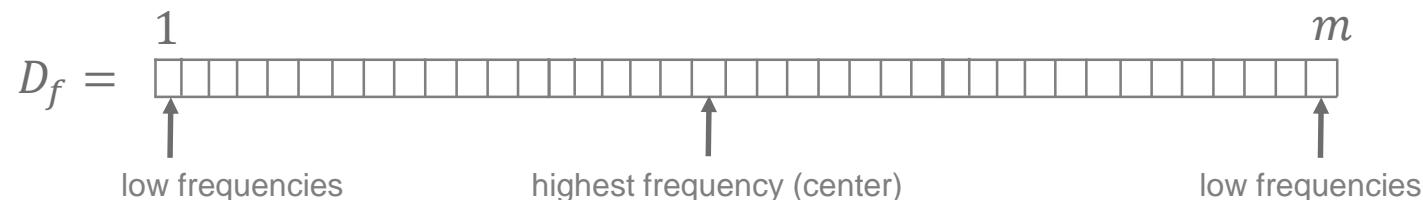
SS 2014

(Course notes for internal use only!)

Comments Ex 4

Generalization of descriptor by using $n = 24$ elements

- Number of elements m in $D = [(x_1 + jy_1), \dots, (x_m + jy_m)]^T$
 - Fourier transform $D_f = fft(D)$

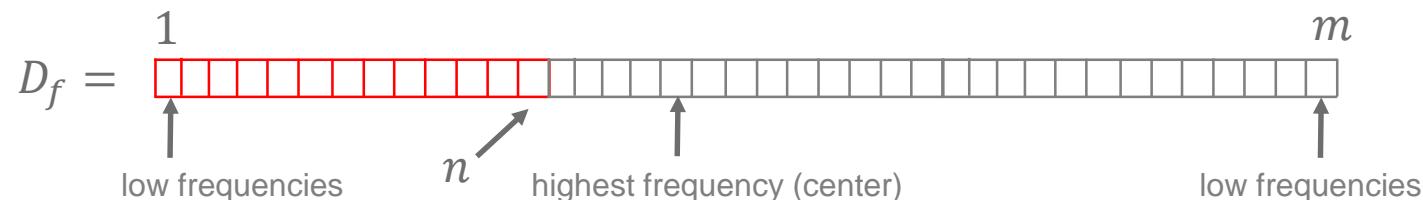


→ Symmetric vector

Comments Ex 4

Generalization of descriptor by using $n = 24$ elements

- Number of elements m in $D = [(x_1 + jy_1), \dots, (x_m + jy_m)]^T$
- Fourier transform $D_f = fft(D)$



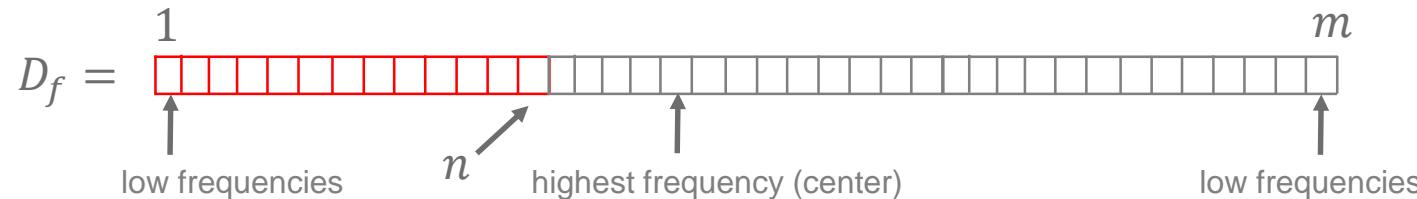
→ Symmetric vector

→ Extract first n values → generalization

Comments Ex 4

Generalization of descriptor by using $n = 24$ elements

- Number of elements m in $D = [(x_1 + jy_1), \dots, (x_m + jy_m)]^T$
 - Fourier transform $D_f = fft(D)$



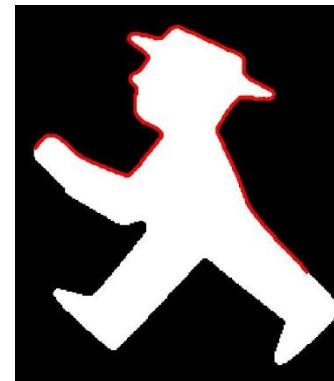
→ Symmetric vector

→ Extract first n values → generalization

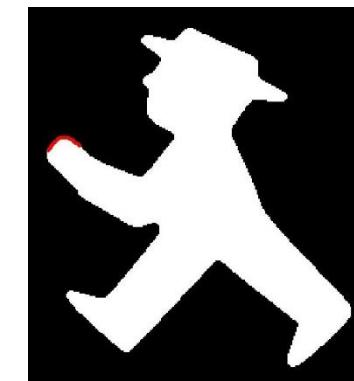
- Fourier transform $D_f = fft(D, n) \rightarrow$ only transforms n elements of D !



n = 842



n = 421



$n = 24$



Exercise 5

SIFT – Scale Invariant Feature Transform

Motivation

- Reliable matching between different views of an object or scene
 - Recognition, Motion tracking, Multiview geometry,...



Motivation

- Reliable matching between different views of an object or scene
 - Recognition, Motion tracking, Multiview geometry,...



Desired Specifications

- Many points
- Repeatable Extraction
- Scale and orientation invariant
- Fast to extract
- Not only positions in image and scale
→ Descriptor vector for each point
- Fast and easy to match

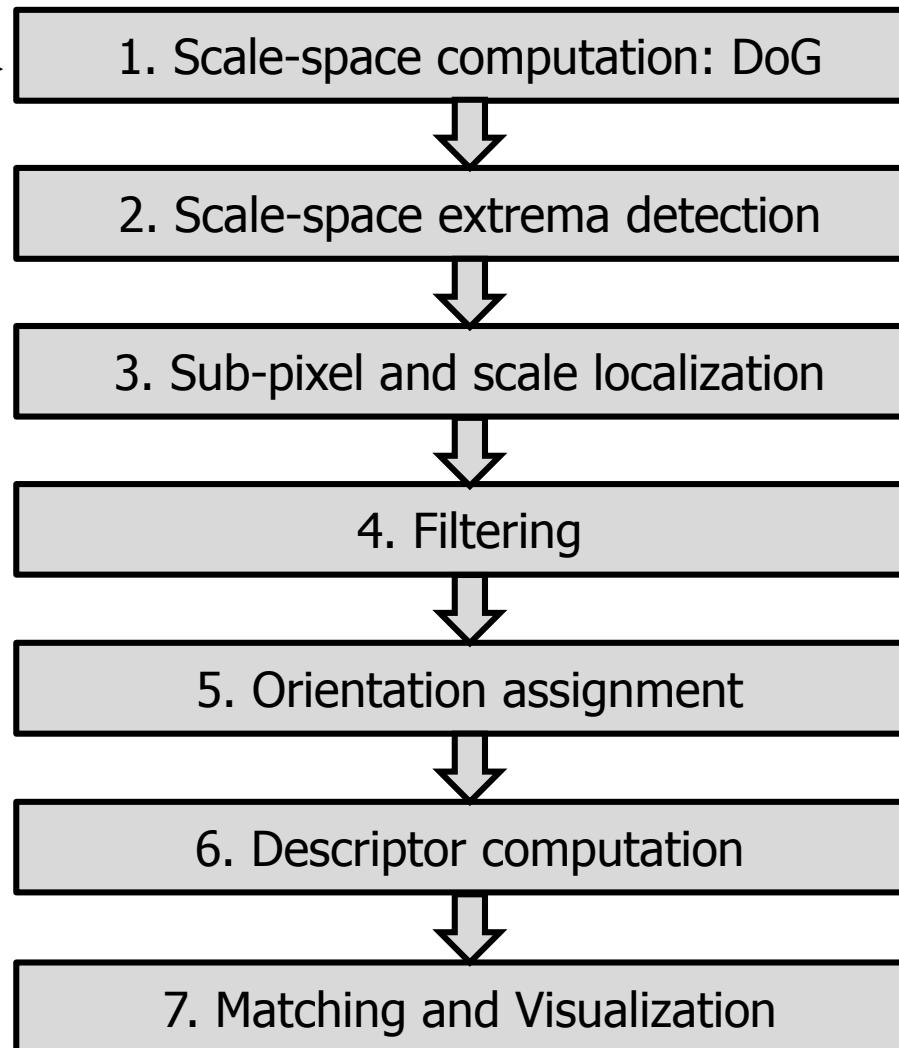
SIFT

- Distinctive image features from scale-invariant keypoints. David G. Lowe, International Journal of Computer Vision, 60, 2 (2004), pp. 91-110.

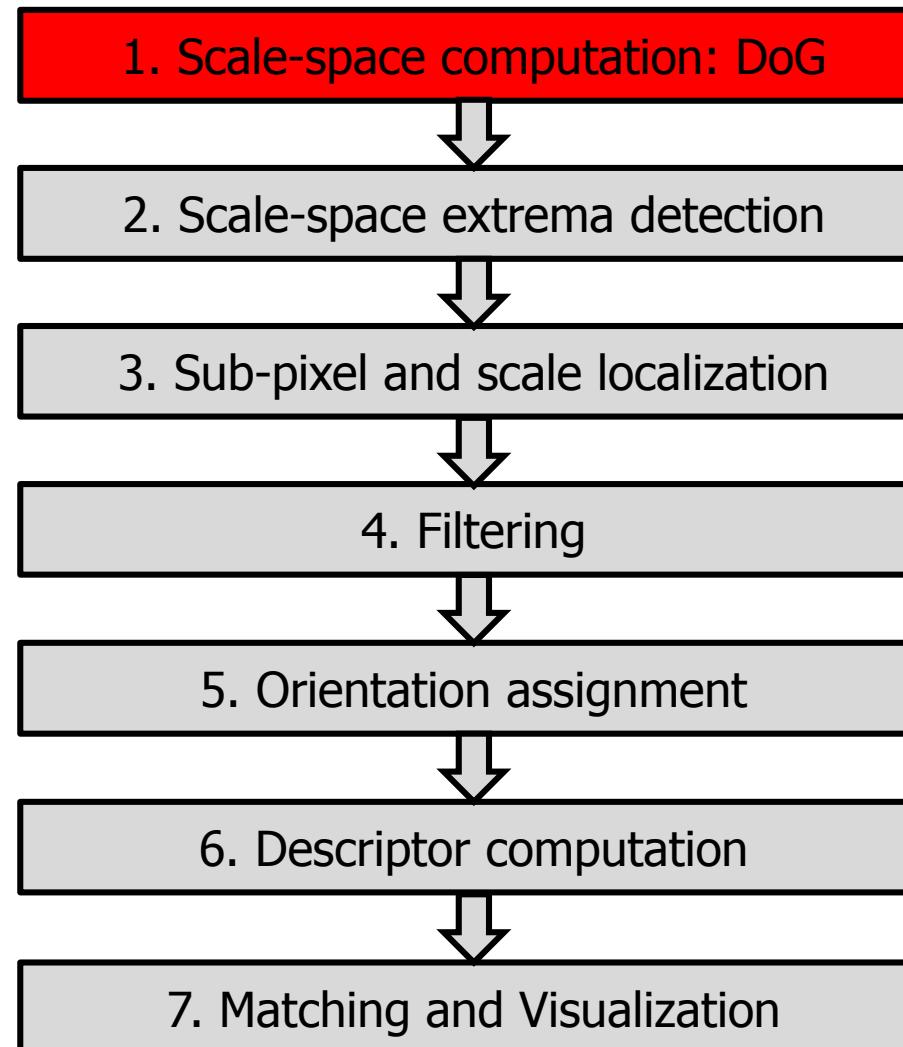
- Finds keypoints
- Generates describing image features for them
- Invariant to image scaling and rotation
- Partially invariant to changes in illumination and camera viewpoints
- Many keypoints are found in typical images
- Highly distinctive

Algorithm Outline

2 Input images →



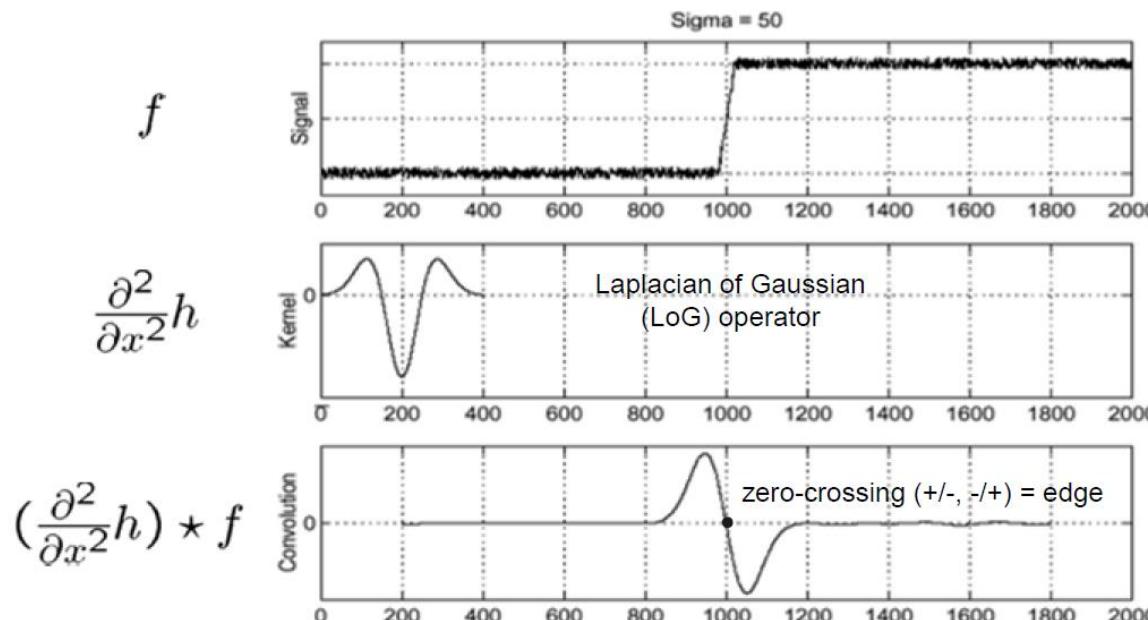
Algorithm Outline



Scale-Space

- Laplacian of Gaussian filtering (LoG)
 - Filtering a signal with second derivatives of Gaussian function
 - Highlights regions of rapid intensity change → zero crossings

Consider: $\frac{\partial^2}{\partial x^2}(h * f)$



Scale-Space

LoG-approximation using Difference of Gaussians (DoG):

- Image blurring using Gaussian filtering:

$$L(x, y, \sigma) = G(x, y, \sigma) * I(x, y)$$

- Difference of Gaussians (DoG):

$$\begin{aligned} D(x, y, \sigma) &= (G(x, y, k\sigma) - G(x, y, \sigma)) * I(x, y) \\ &= L(x, y, k\sigma) - L(x, y, \sigma) \end{aligned}$$

- Easy and efficient to compute

Scale-Space: Example

- Input image



Scale-Space: Example

Blurred image



DoG-Image



$$\sigma = 2.0$$

Scale-Space: Example

Blurred image



DoG-Image



$$\sigma = 4.0$$

Scale-Space: Example

Blurred image



DoG-Image



$$\sigma = 8.0$$

Scale-Space: Example

Blurred image



DoG-Image



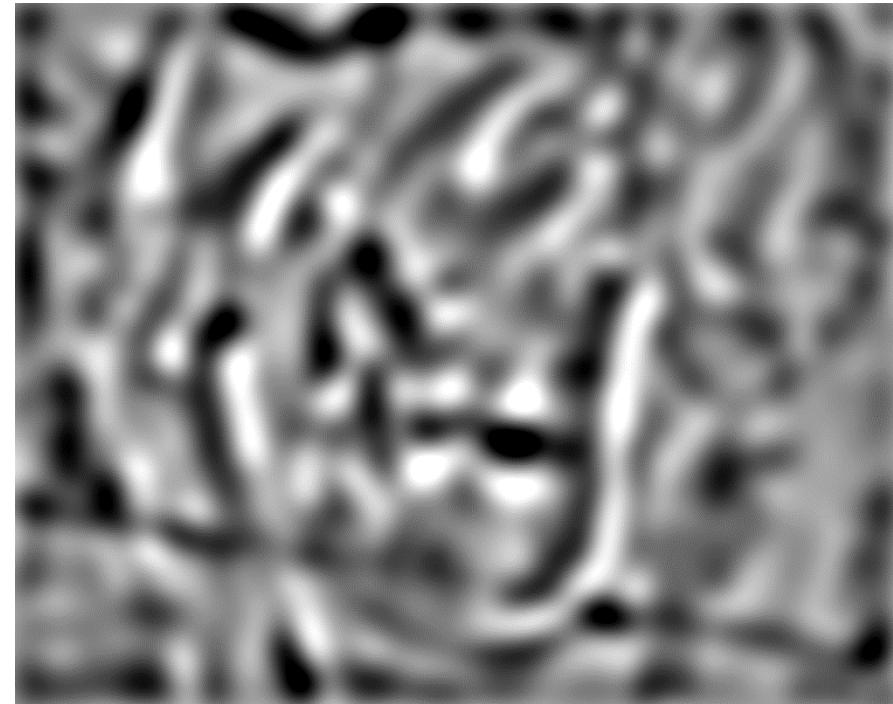
$$\sigma = 16.0$$

Scale-Space: Example

Blurred image



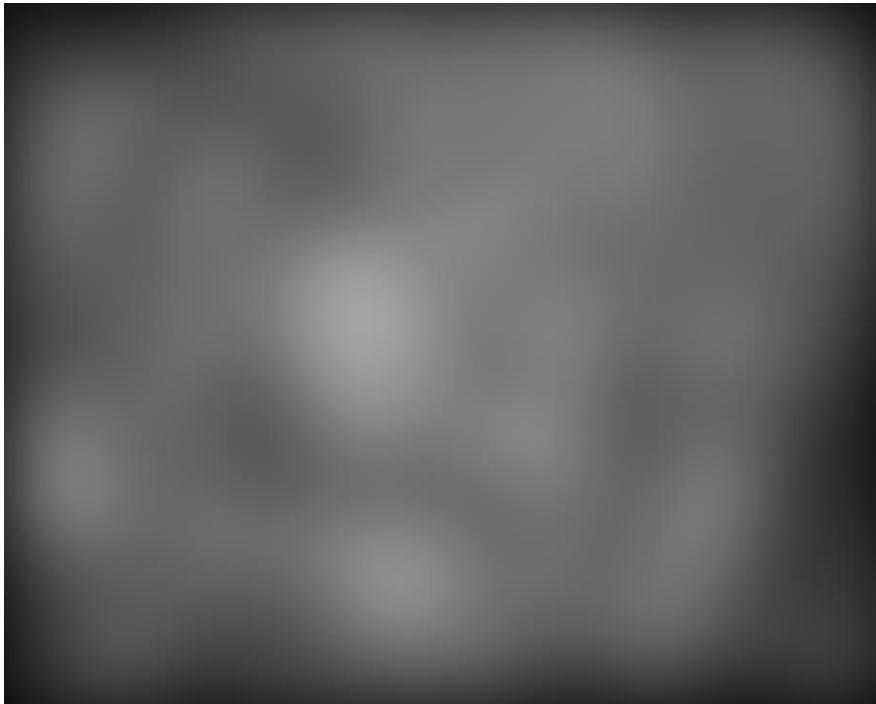
DoG-Image



$$\sigma = 32.0$$

Scale-Space: Example

Blurred image



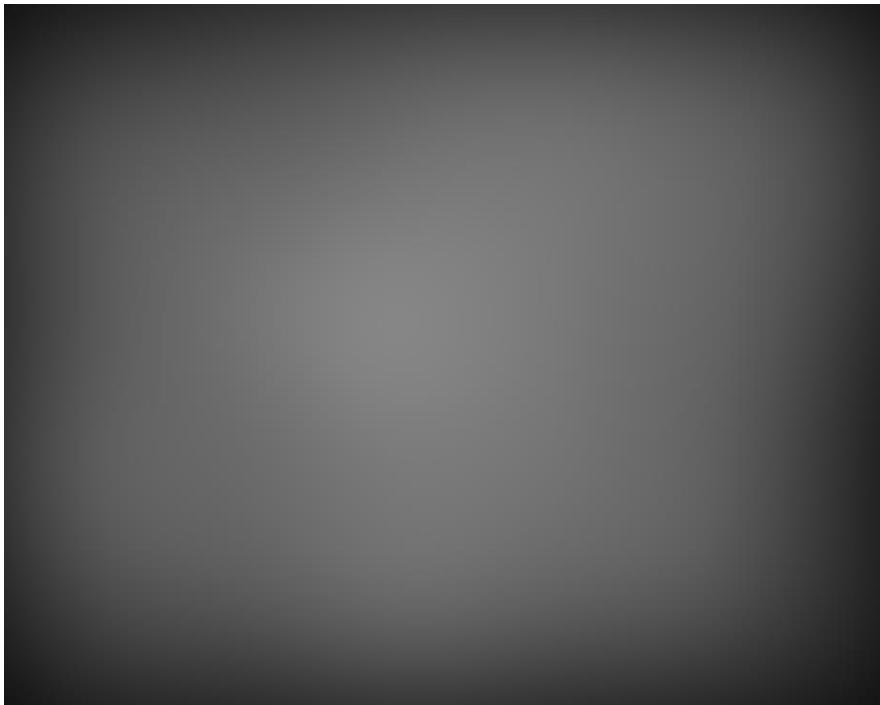
DoG-Image



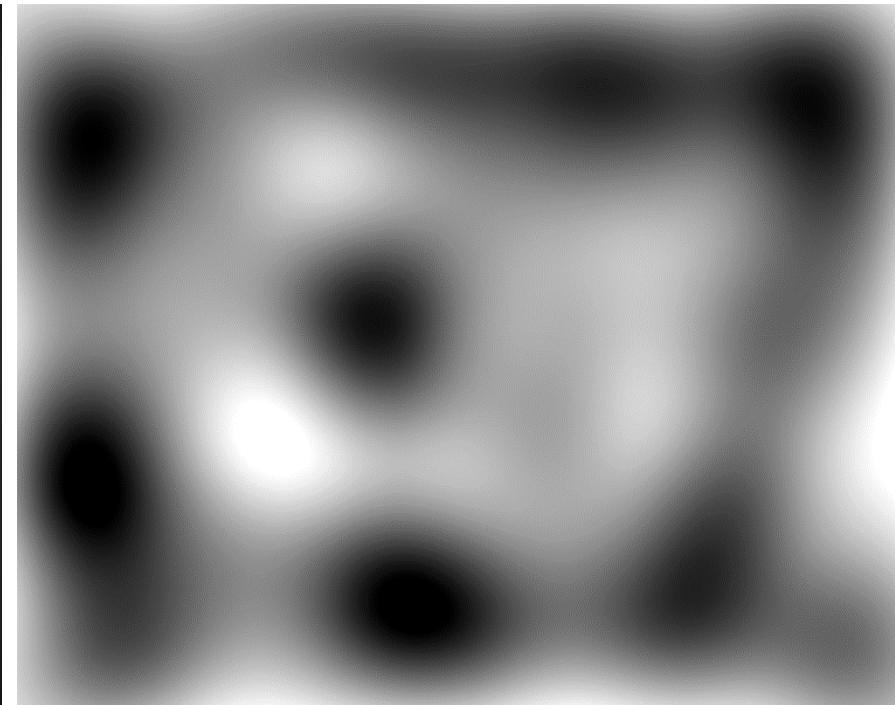
$$\sigma = 64.0$$

Scale-Space: Example

Blurred image



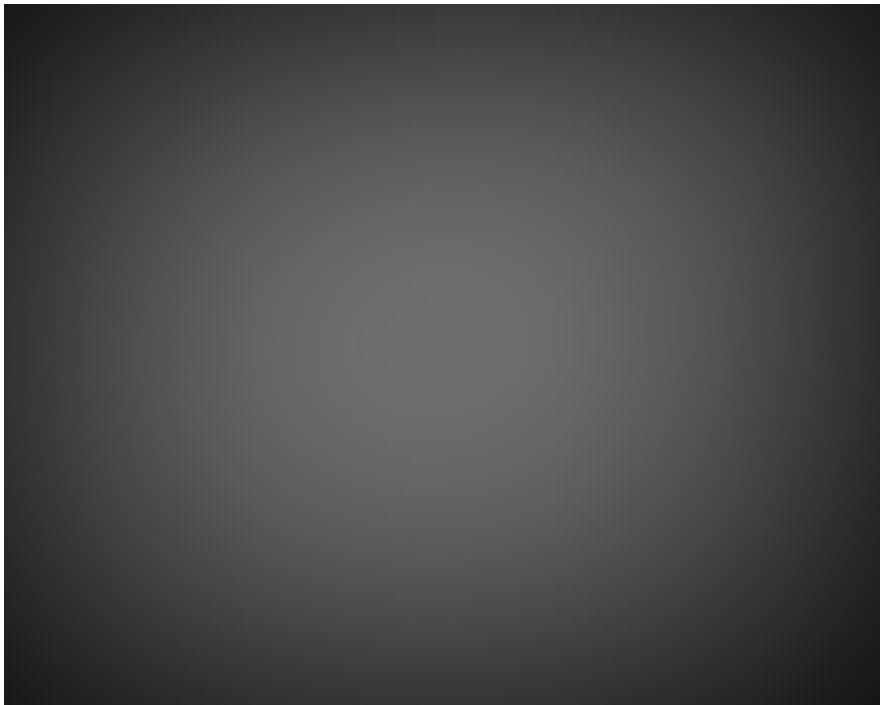
DoG-Image



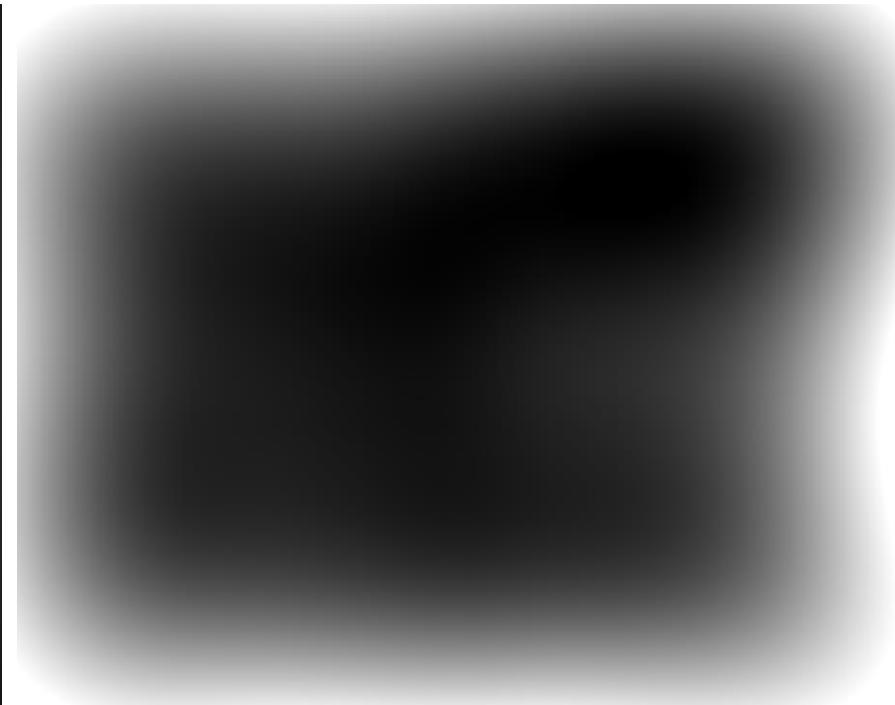
$$\sigma = 128.0$$

Scale-Space: Example

Blurred image



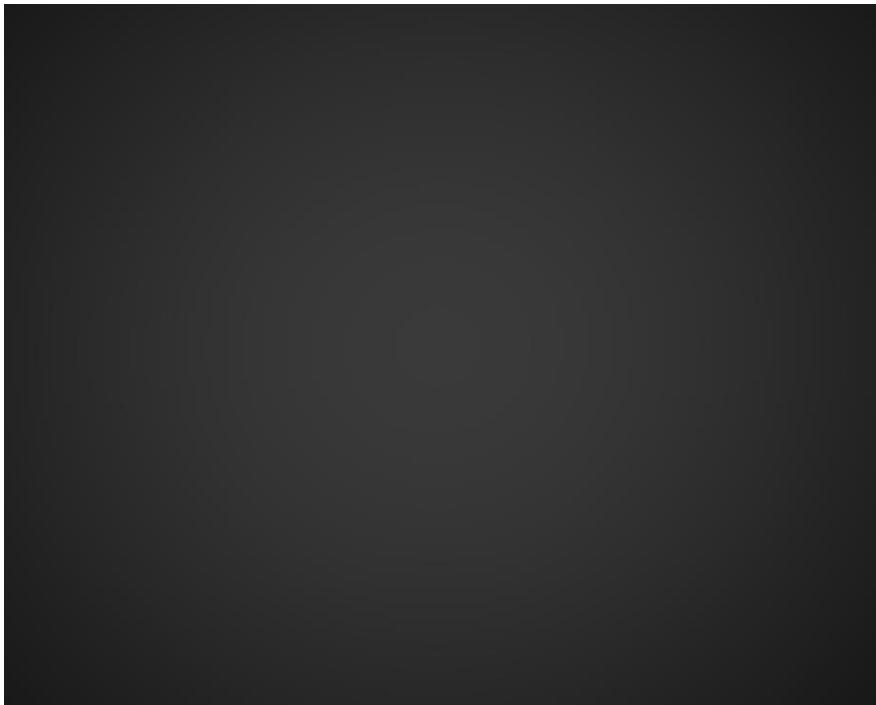
DoG-Image



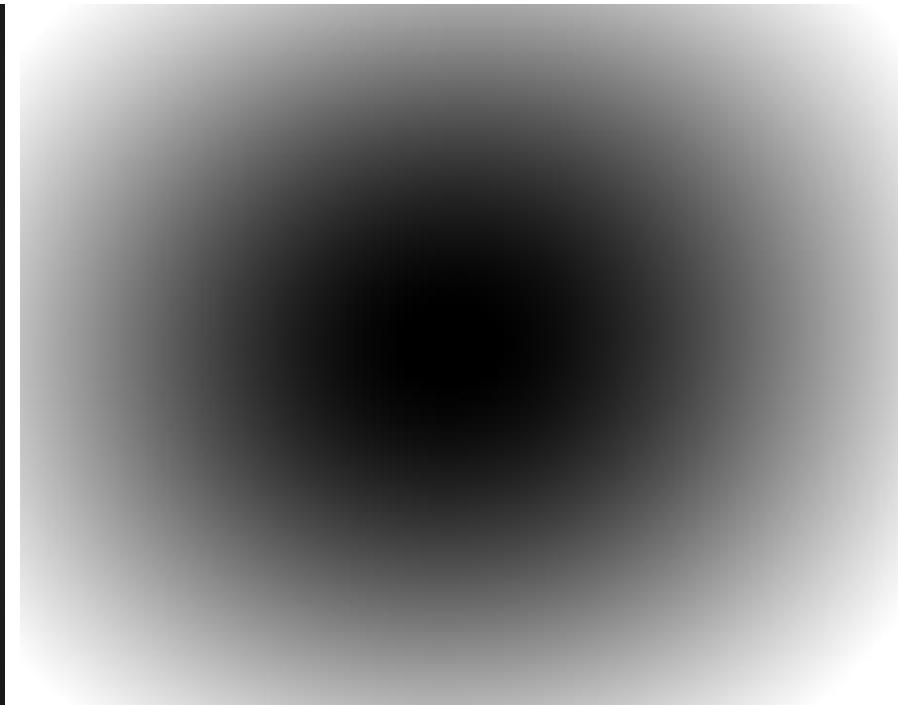
$$\sigma = 256.0$$

Scale-Space: Example

Blurred image



DoG-Image



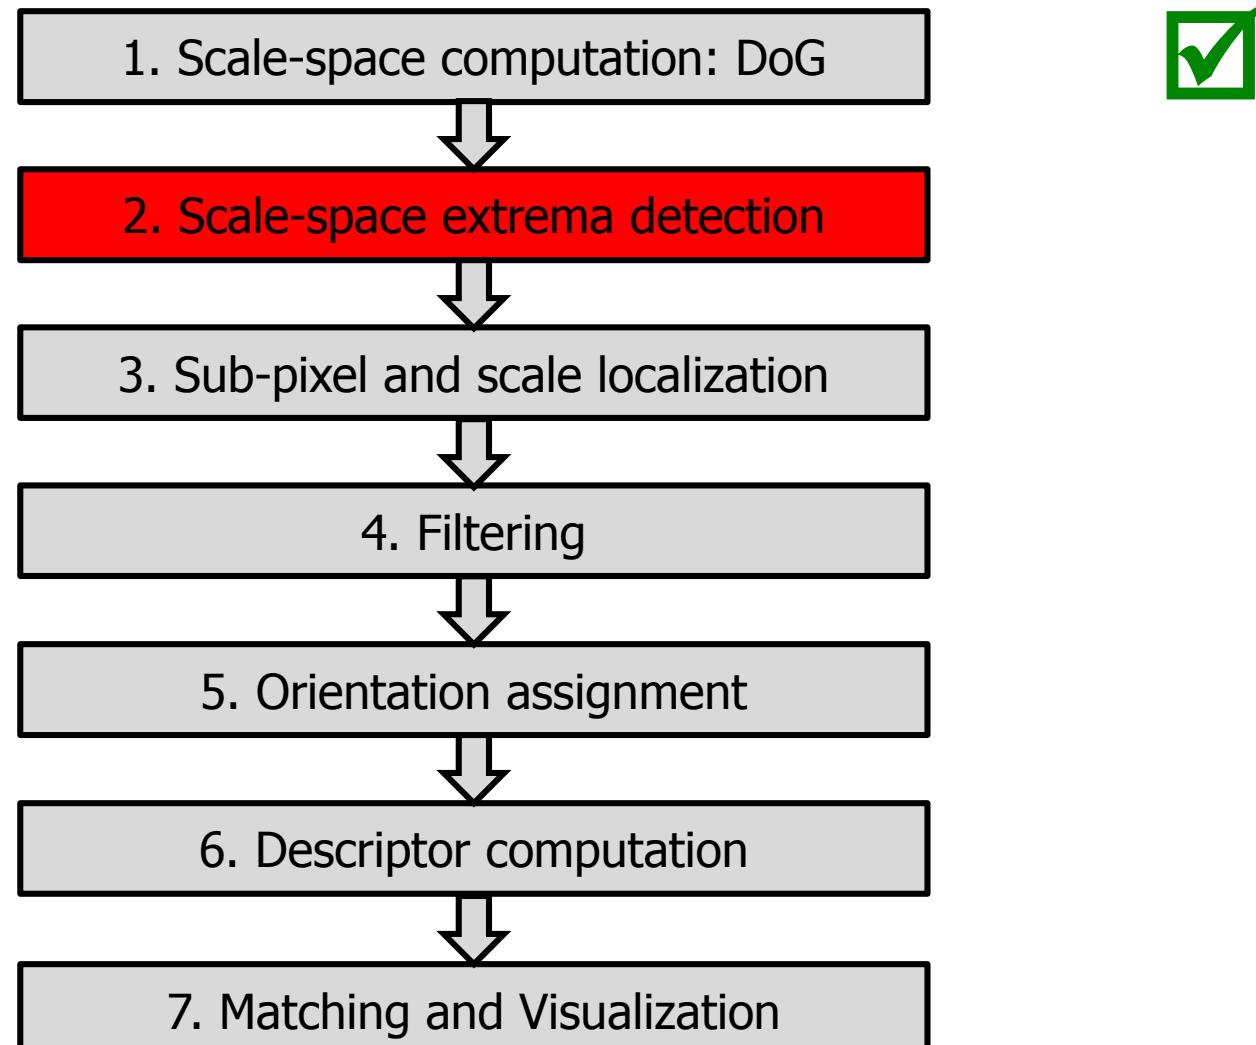
$$\sigma = 512.0$$

Scale-Space: Prior Smoothing

Prior Gaussian smoothing of the input image

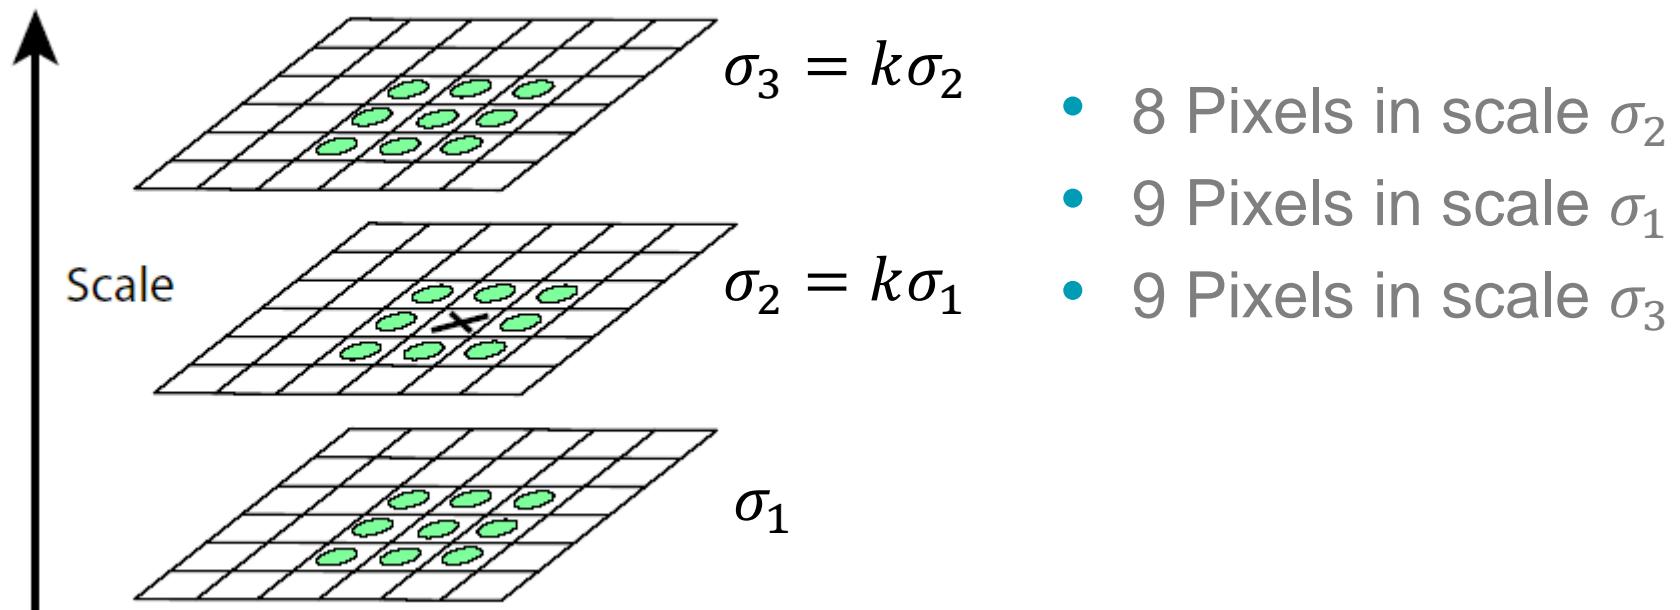
- Increasing σ also increases **robustness** and **repeatability**
- $\sigma = 1.6$ has shown to be good (found empirically)
- Smoothing discards highest frequencies
→ Doubling the image initially increases number keypoints

Algorithm Outline

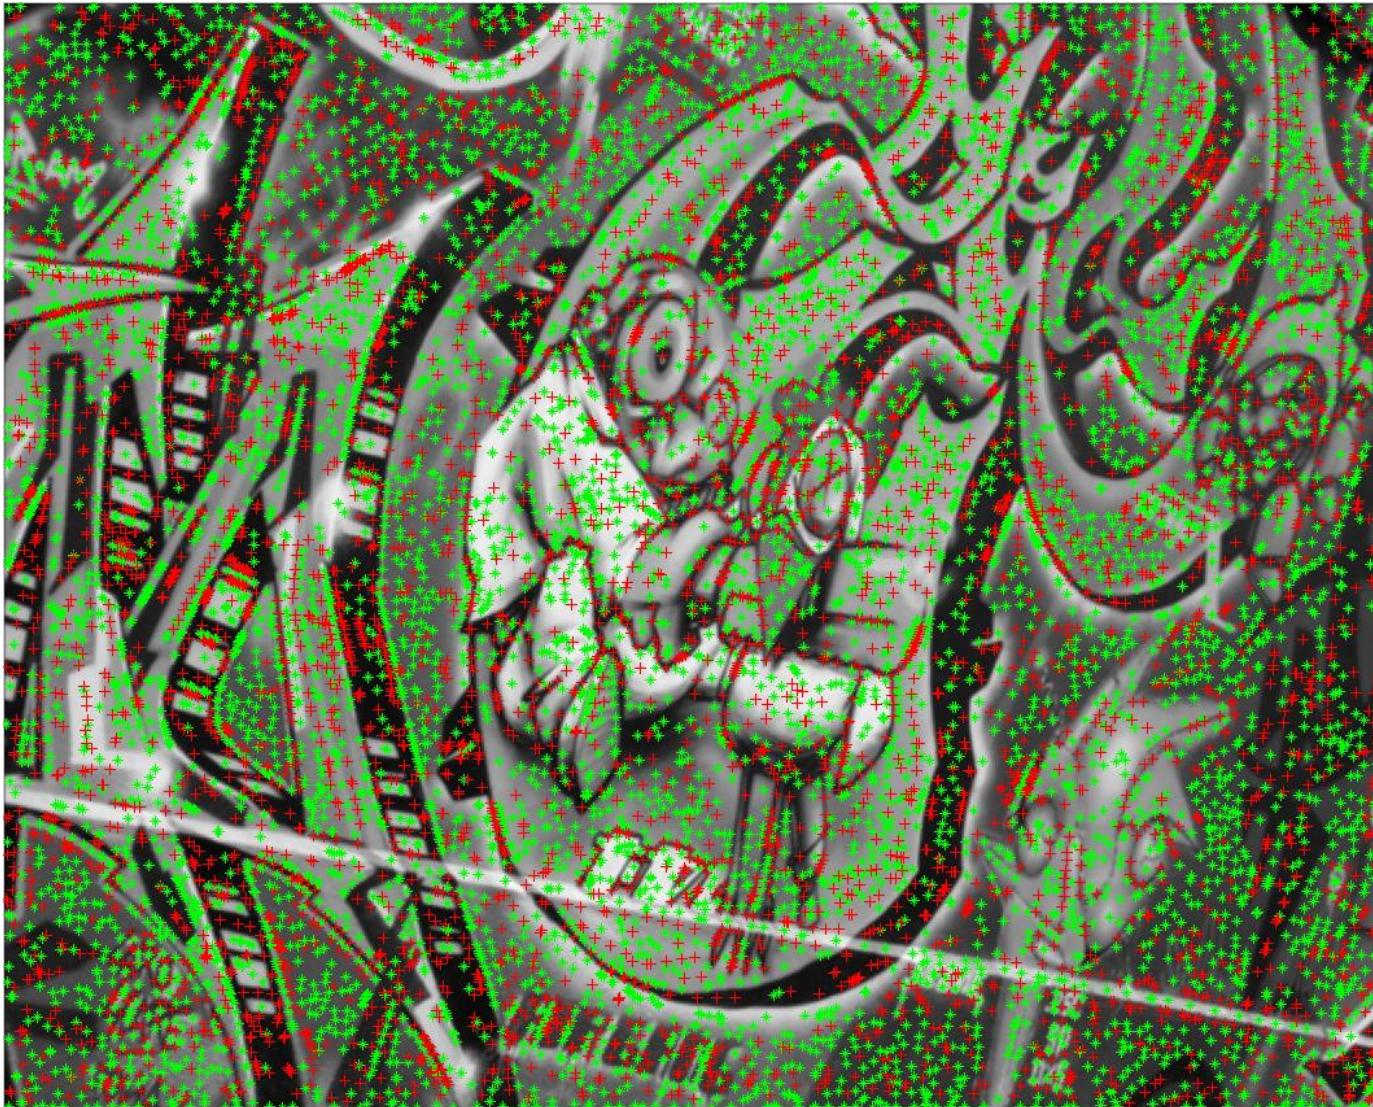


Extrema Detection

- Find local Maxima and Minima in DoG-Pyramid
→ Check local neighborhood

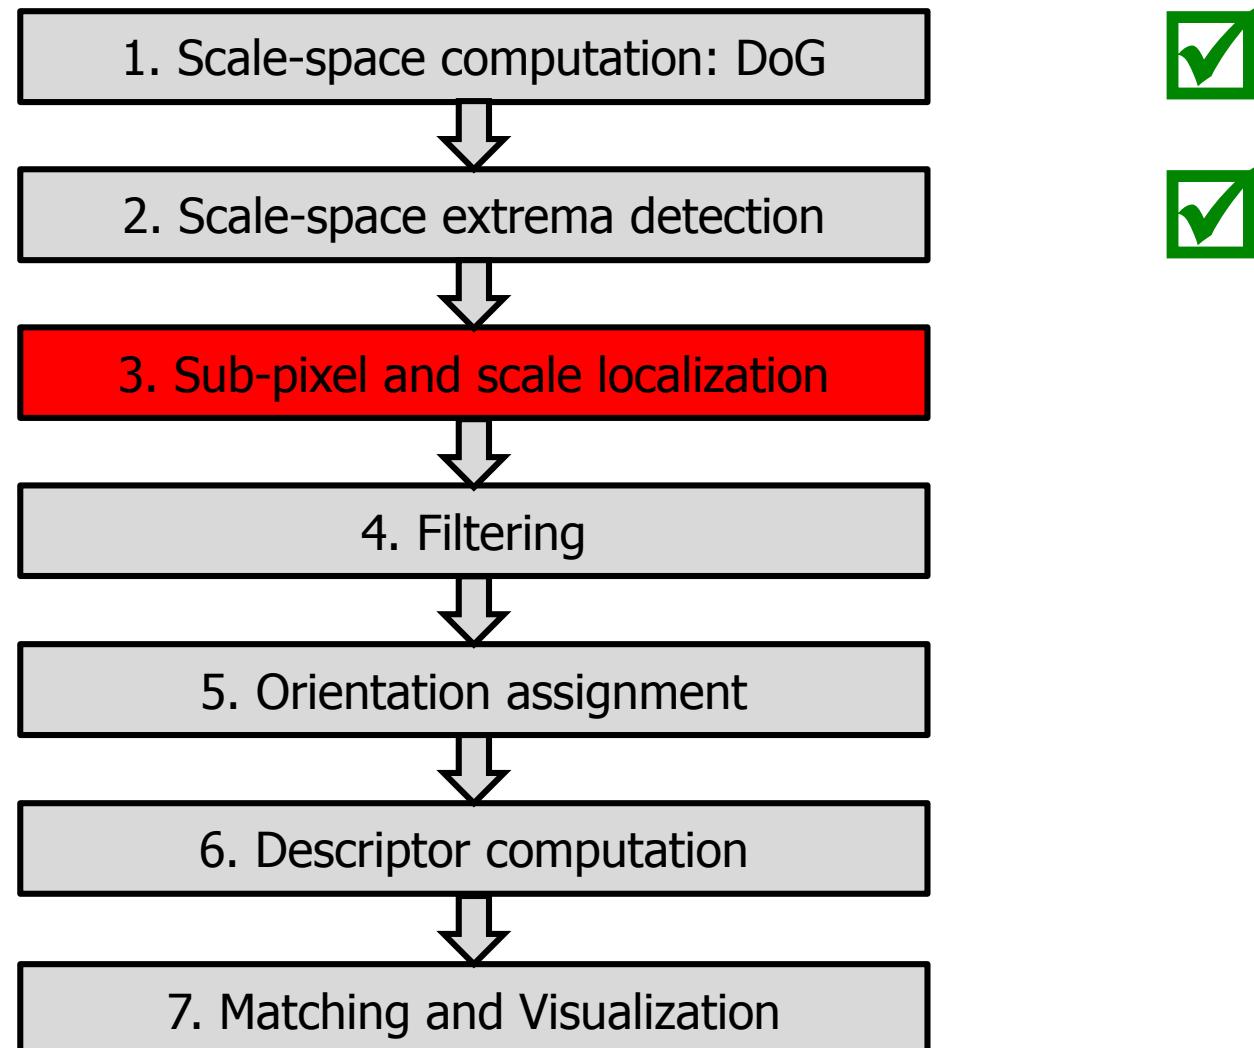


Extrema Detection Result



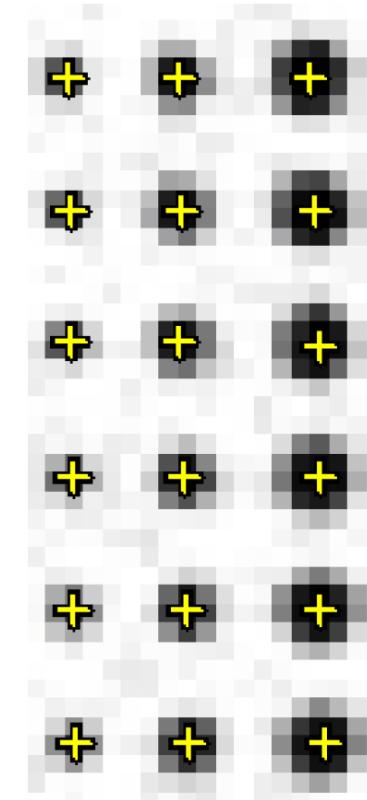
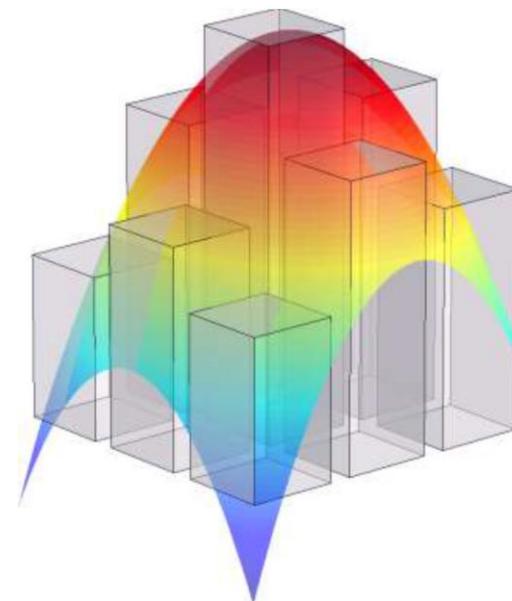
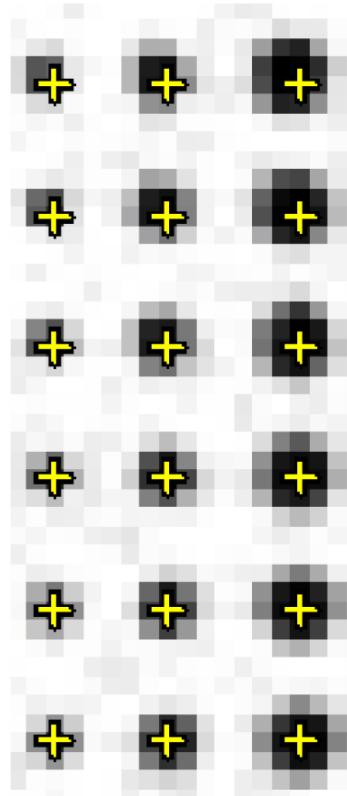
Red: minima
Green: maxima

Algorithm Outline



Localization

- Aim: Find the subpixel position of keypoint candidates in position and scale (3D)
- Example for 2D: Paraboloid fitting



Localization

- Fit 3D quadratic functions to local keypoint candidates

- Taylor expansion with keypoint candidate X as origin

$$D(X) = D + \frac{\partial D^T}{\partial X} X + \frac{1}{2} X^T \frac{\partial^2 D}{\partial X^2} X, \text{ where } X = (x, y, \sigma)$$

- Computation of derivatives: simple pixel differences

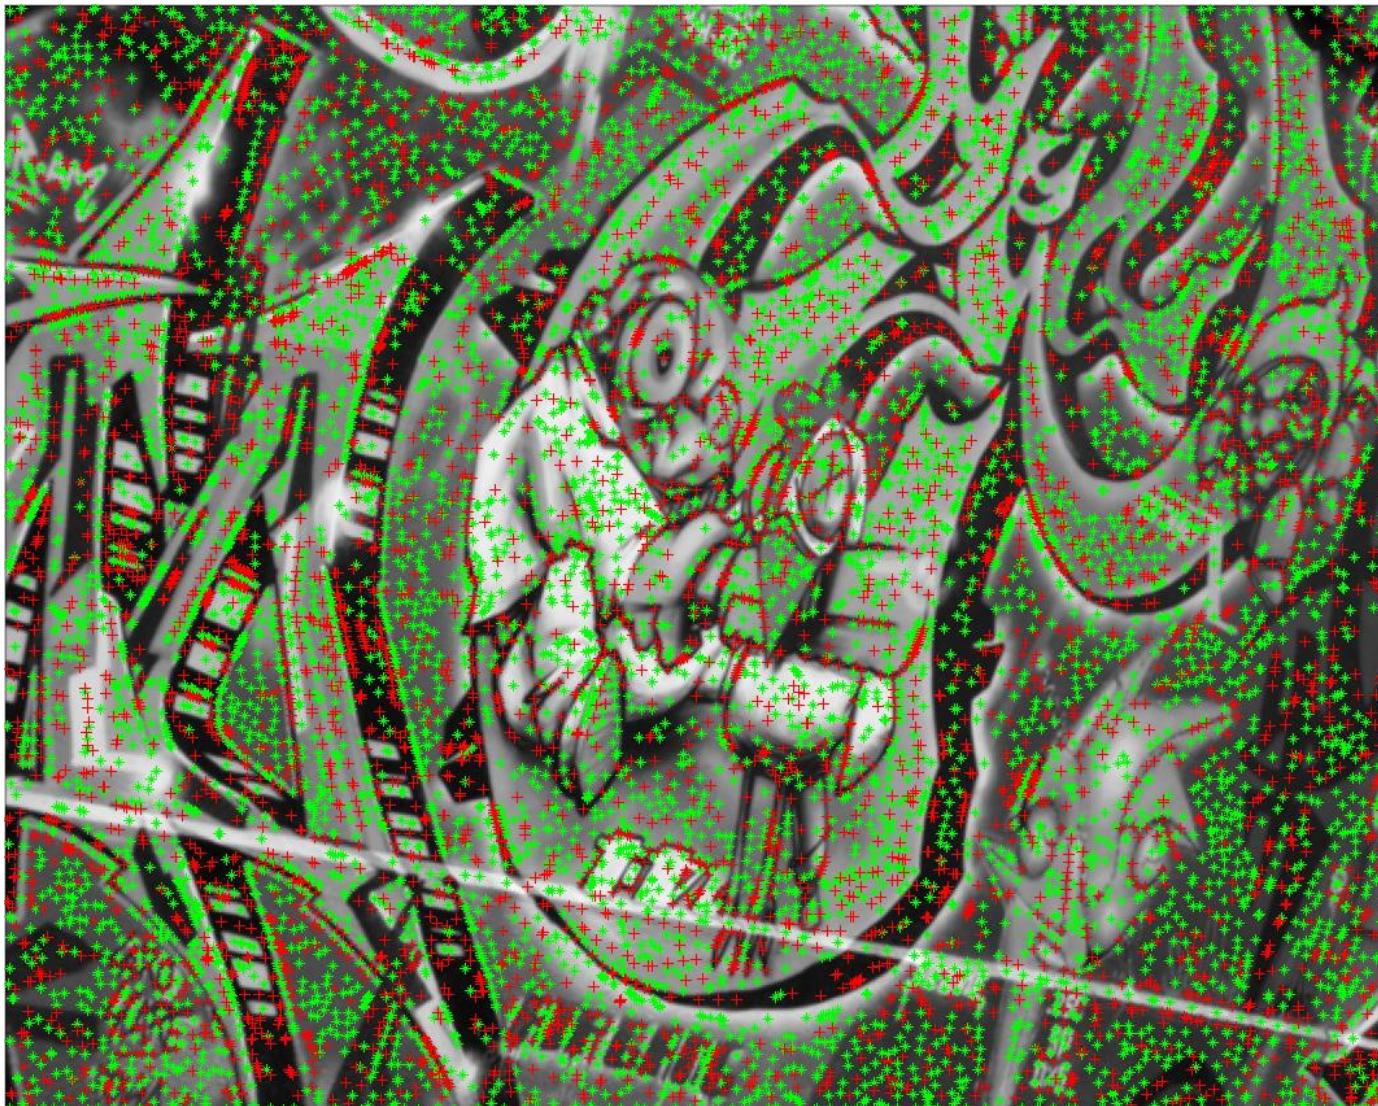
- Set first derivative wrt. X to zero

$$0 = \frac{\partial D}{\partial X} + \frac{\partial^2 D}{\partial X^2} \hat{X}$$

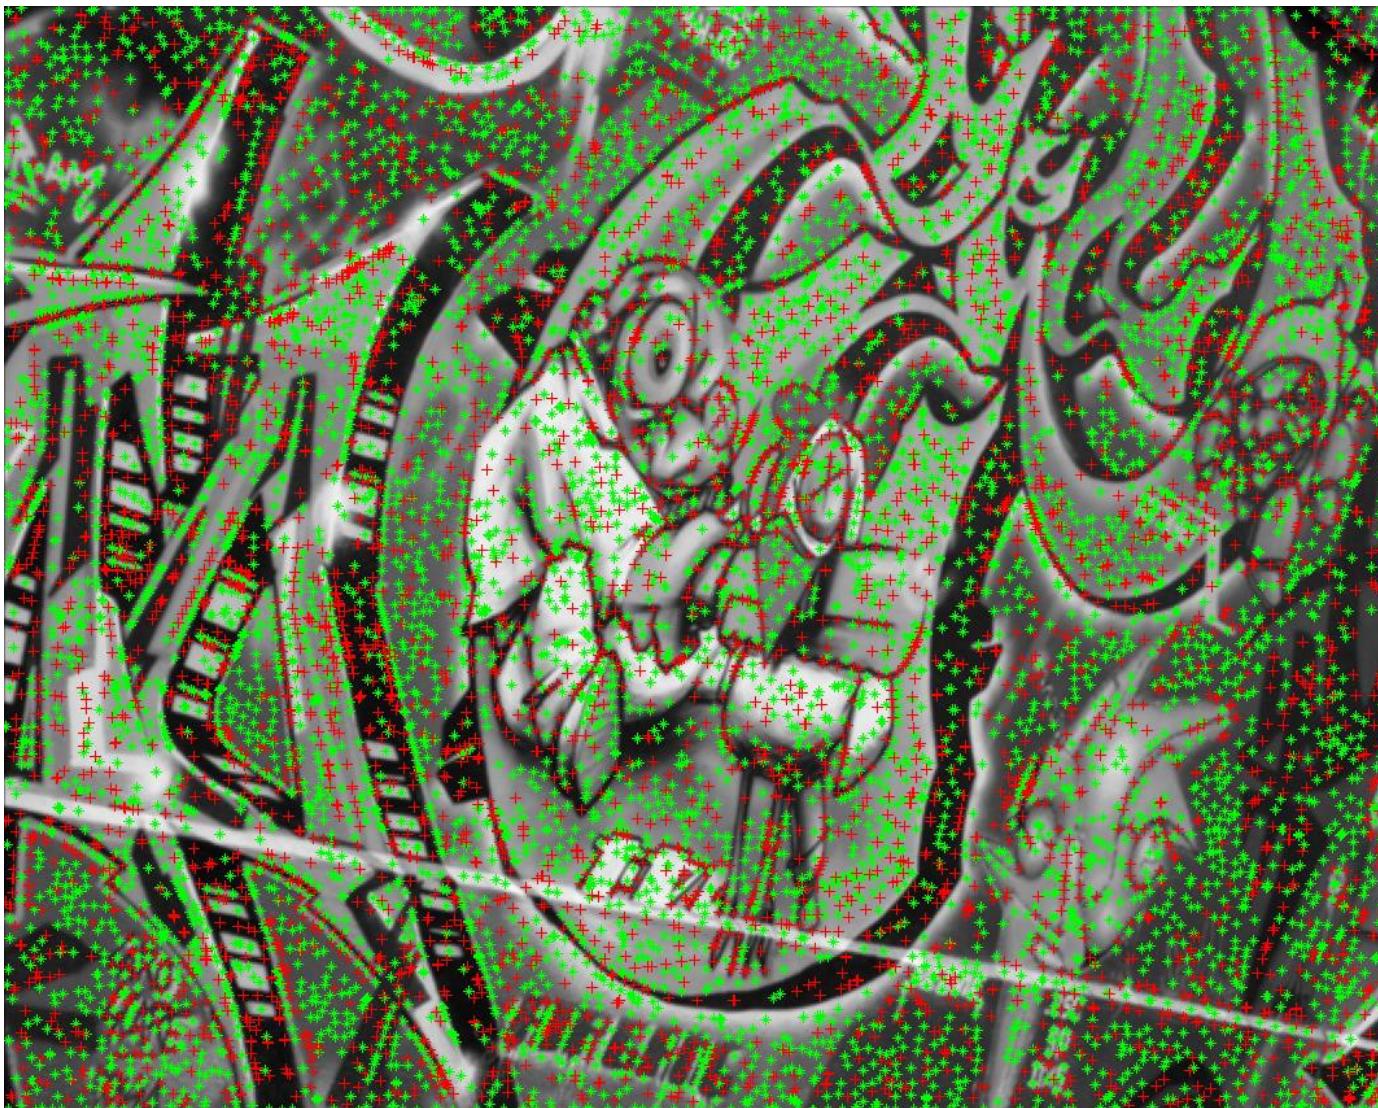
- Solve corresponding linear 3×3 system

$$\hat{X} = - \frac{\partial^2 D^{-1}}{\partial X^2} \frac{\partial D}{\partial X}$$

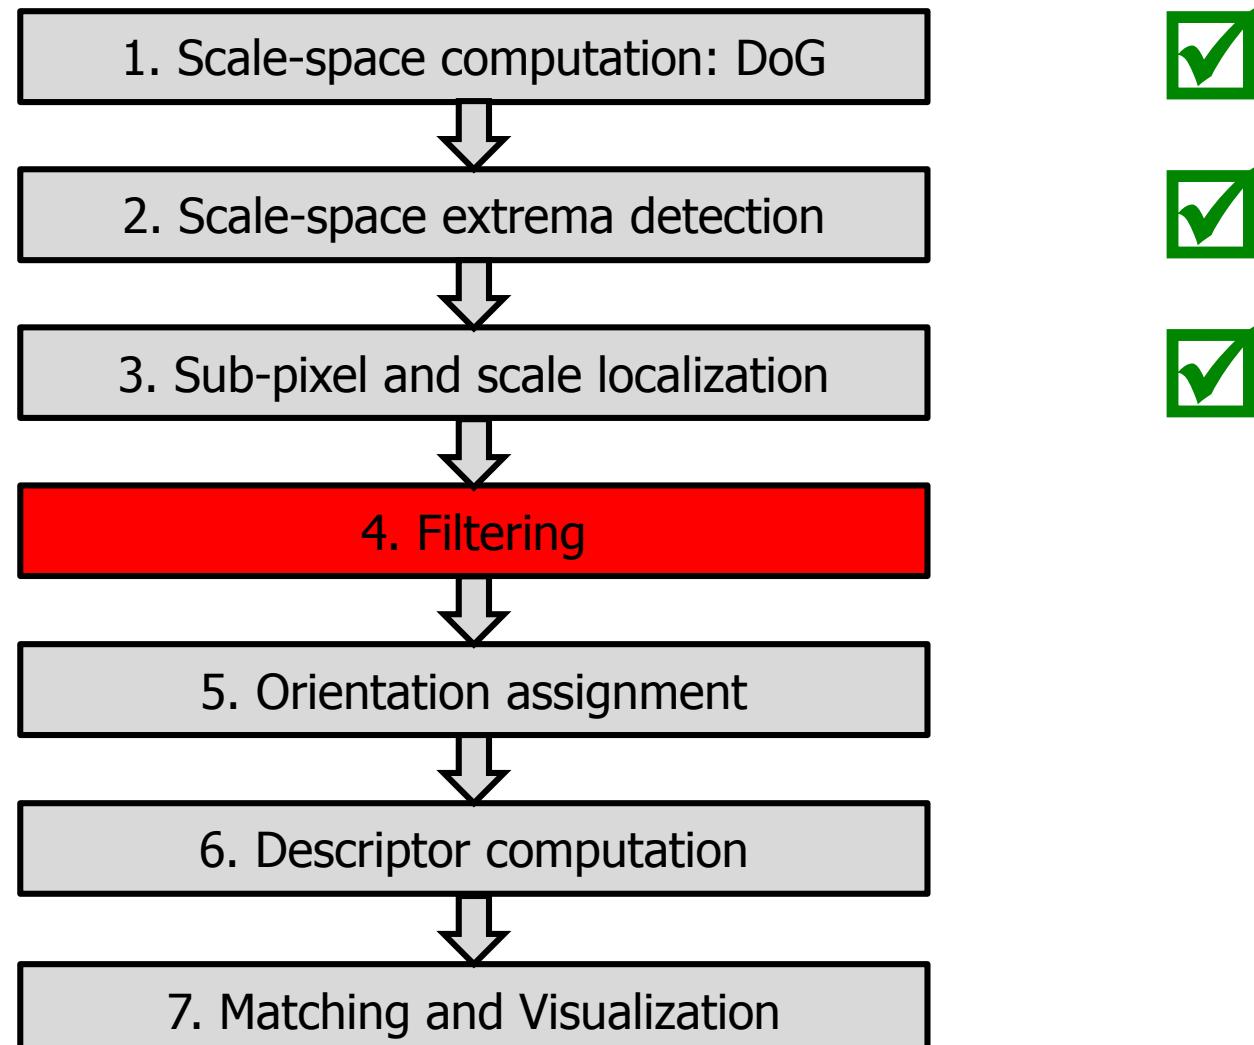
Localization: Initial Points



Localization: Estimated Locations



Algorithm Outline



Filtering

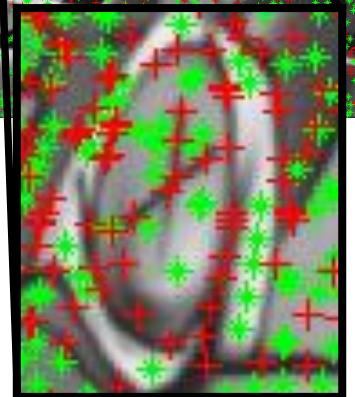
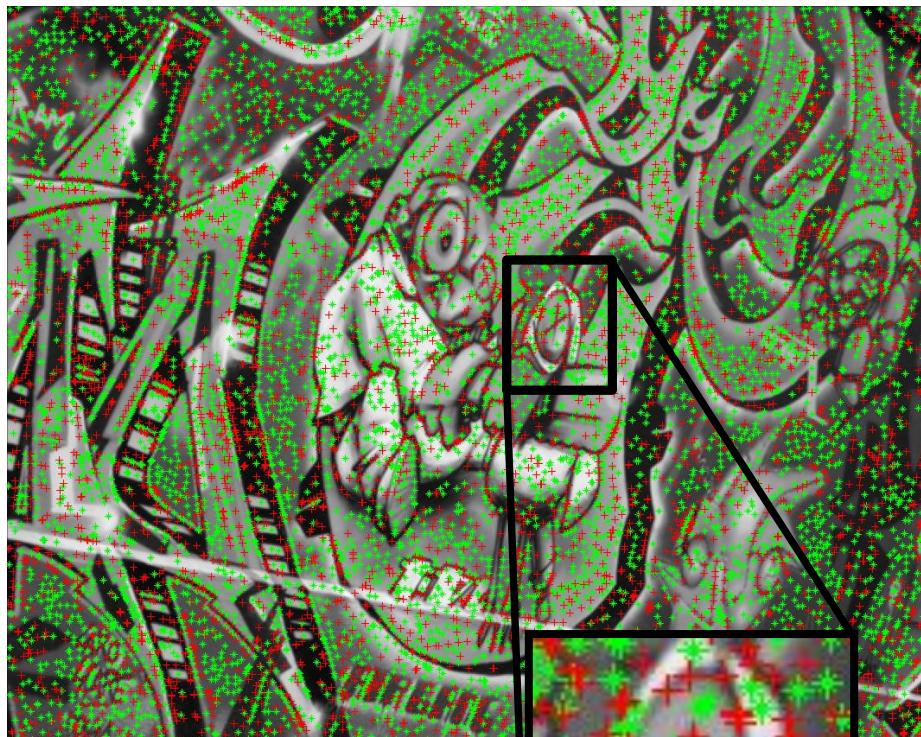
- Filter low contrast points, if

$$|D(\hat{X})| = \left| D + \frac{1}{2} \frac{\partial D^T}{\partial X} \hat{X} \right| < 0.03$$

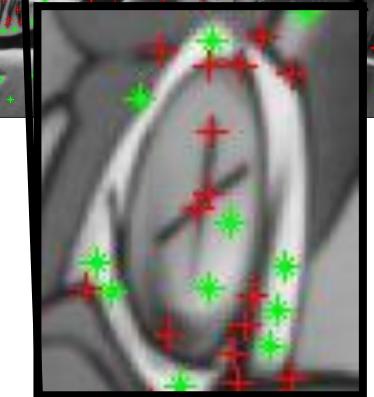
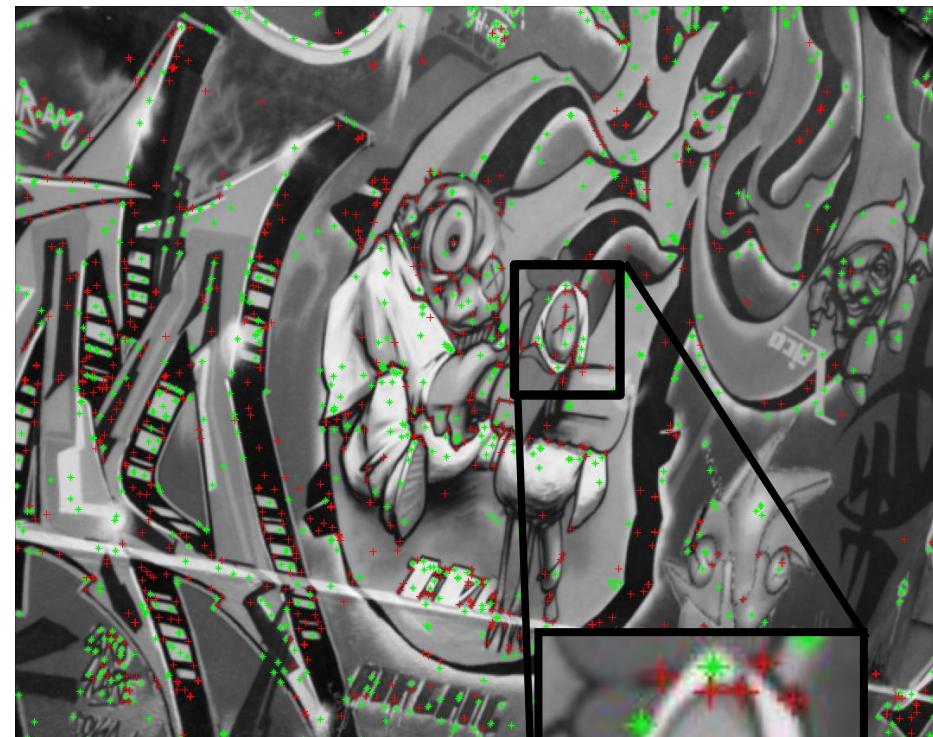
- More stability: Eliminate edge-responses
 - One large principal curvature across the edge but a small one in perpendicular direction
 - Principal curvatures from Hessian: $H = \begin{bmatrix} D_{xx} & D_{xy} \\ D_{xy} & D_{yy} \end{bmatrix}$
 - If $\frac{\text{Trace}(H)}{\text{Det}(H)} > \frac{(r+1)^2}{r}$ → throw point away (e.g. with $r = 10$)

Filtering Results

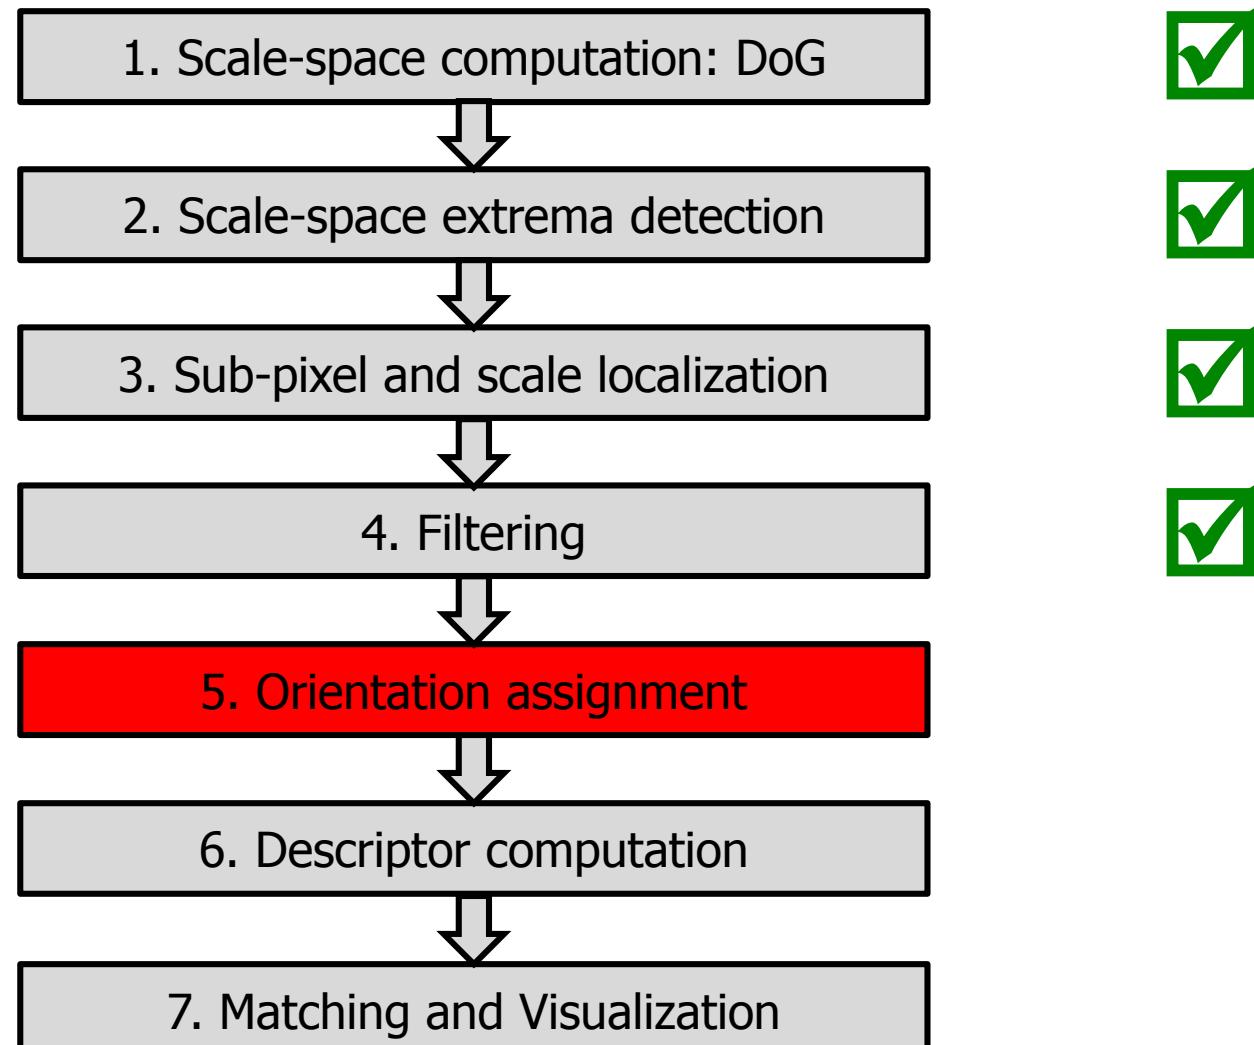
Localized Points



Filtered Points



Algorithm Outline



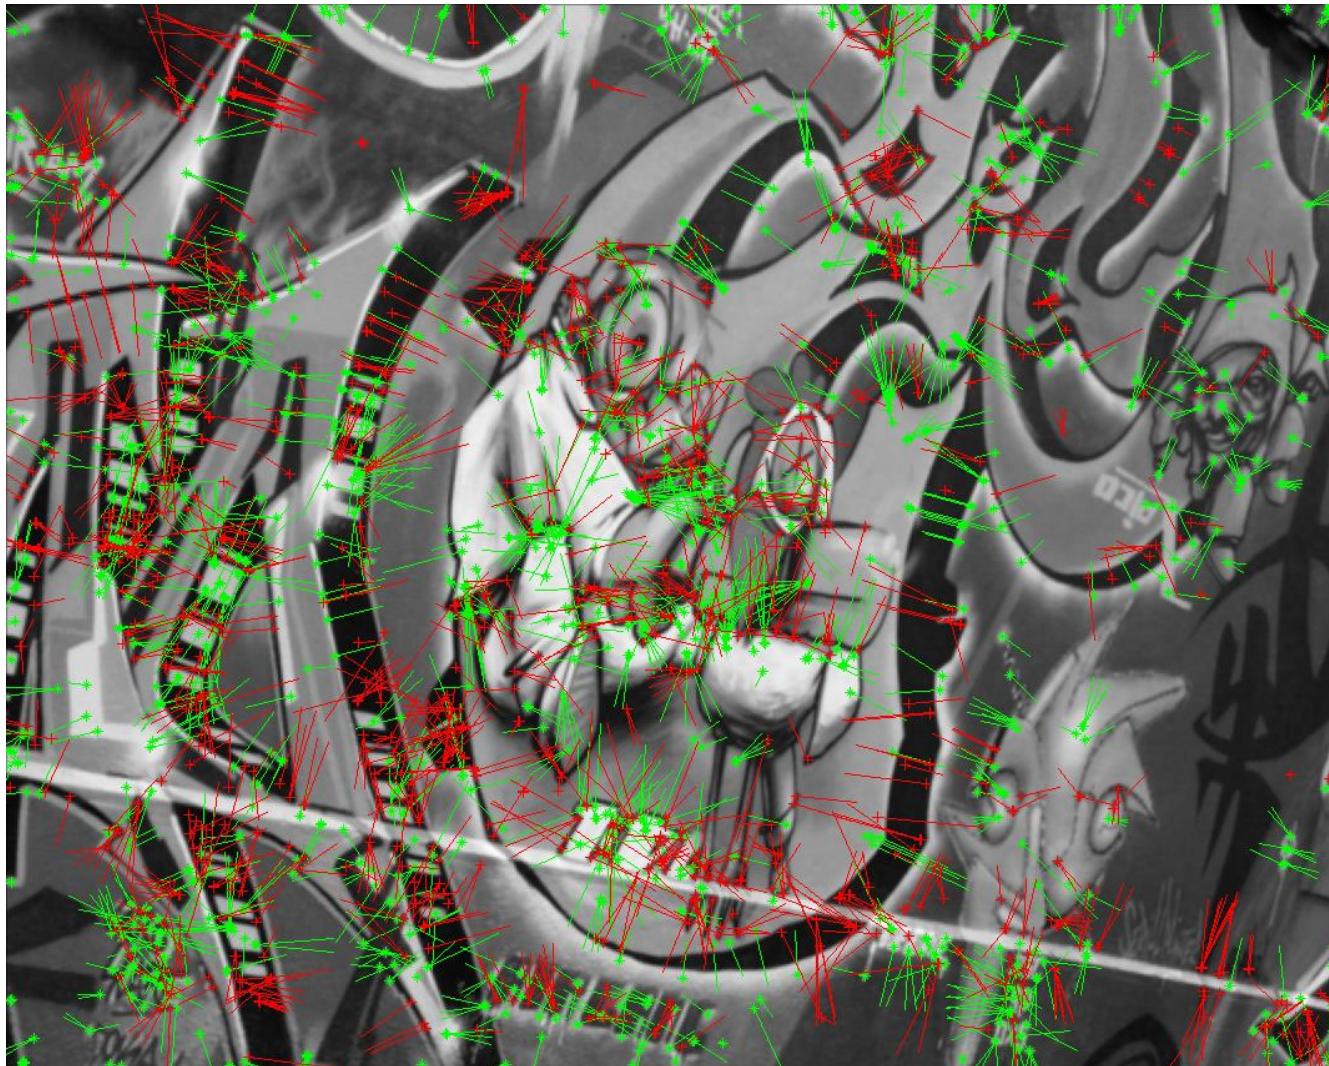
Assign Gradient Orientations

- **Idea:** achieving **orientation invariance** by computing descriptors relative to keypoints main orientation
- Computation of Gradient orientations θ and magnitudes m for all scales and pixels using Gaussian filtered images L

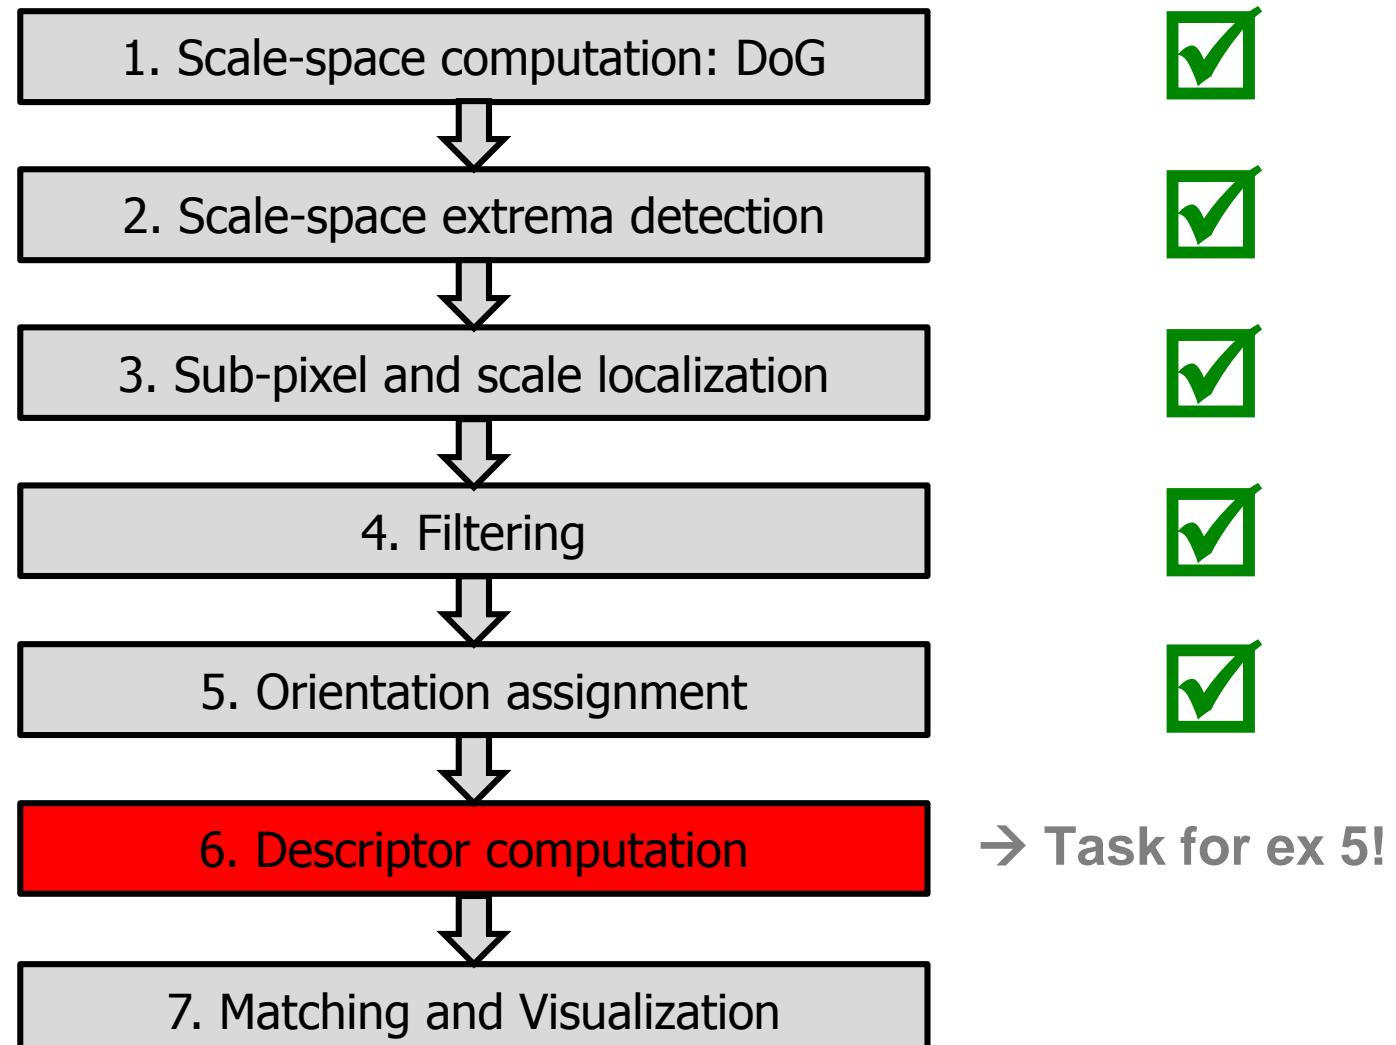
$$m(x, y) = \sqrt{(L(x + 1, y) - L(x - 1, y))^2 + (L(x, y + 1) - L(x, y - 1))^2}$$
$$\theta(x, y) = \text{atan} \left((L(x, y + 1) - L(x, y - 1)) / (L(x + 1, y) - L(x - 1, y)) \right)$$

- Use this information to assign main orientation of local neighborhood to each keypoint

Orientation Assignment: Result



Algorithm Outline



Descriptor Computation: Overview

- **Inputs**

- Gradient magnitude and angle images for each scale
- Keypoint positions + their main local orientation

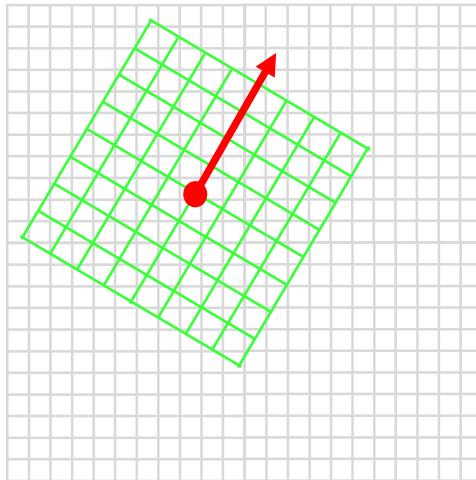
- **Outputs**

- 128-element descriptor vector for each keypoint
 - This has to be computed!
- 5-element vector for each keypoint: $[y, x, \sigma, \theta_m, m_m]$, where
 - (y, x) : subpixel image coordinates
 - σ : scale
 - θ_m : main gradient orientation
 - m_m : main gradient magnitude

→ These values are already available, rearrange them (for visualization)

Descriptor Computation: Overview

Keypoint and gradient direction

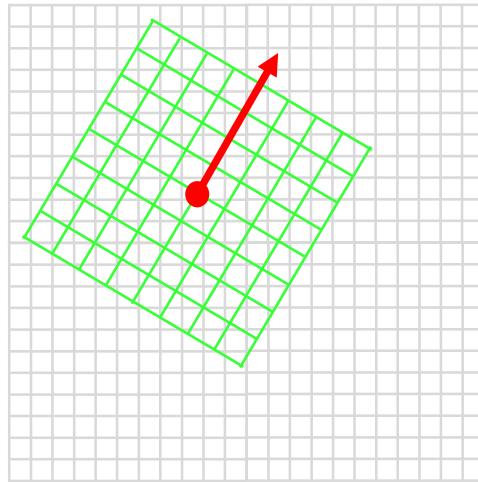


Main idea: Use gradient magnitude and orientation of local neighborhood to compute **scale and rotation invariant** descriptor

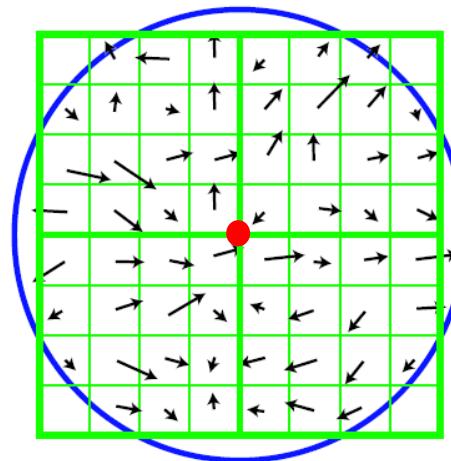
- Rotate sampling window by assigned main gradient direction

Descriptor Computation: Overview

Keypoint and gradient direction



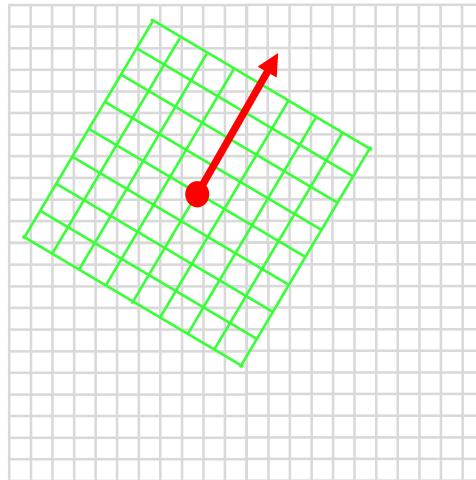
Sampled gradients



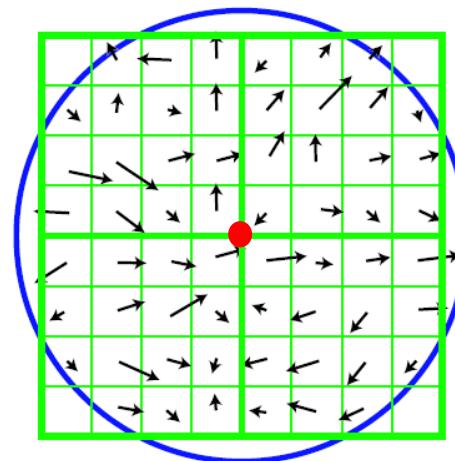
- Rotate sampling window by assigned main gradient direction
- Sample gradient magnitudes and orientations and weight magnitudes using a Gaussian function (blue circle)

Descriptor Computation: Overview

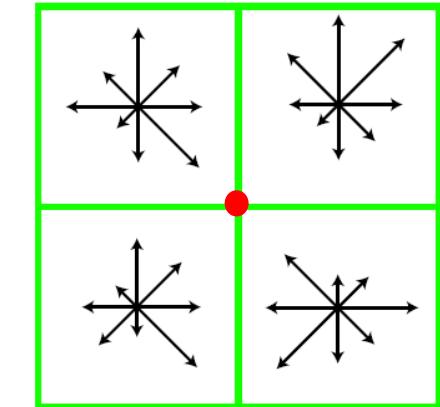
Keypoint and gradient direction



Sampled gradients



Descriptor



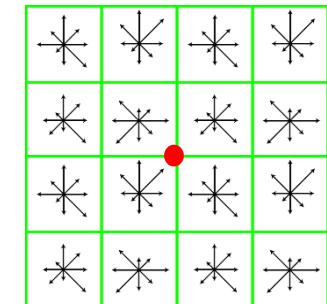
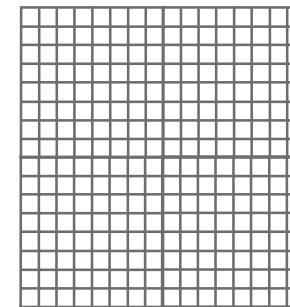
- Rotate sampling window by assigned main gradient direction
- Sample gradient magnitudes and orientations and weight magnitudes using a Gaussian function (blue circle)
- Build descriptor from orientation histograms of gradient magnitudes

Descriptor Computation: Some Facts

1) Window size w for descriptor: 16×16 px

→ 16 histograms

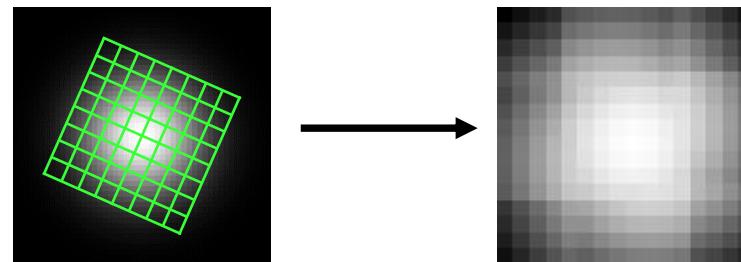
→ Each histogram is build using 4×4 px



2) Gaussian weighting window: $\sigma = 0.5 \cdot w = 8 \rightarrow 49 \times 49$ filter

→ Filter has to be resampled to size of local window (16×16)

→ Resampling needed wrt position of keypoint



3) Use Gradient magnitudes and orientations from the nearest scale σ of keypoints

Descriptor Computation: Algorithm

- compute the 49×49 Gaussian filter ($\sigma = 8$)
- **for all scales**
 - **for all points in current scale**
 - Initialize arrays x_c, y_c with local image coordinates for a 16×16 neighborhood
 - Rotate local coordinates x_c and y_c using main orientation θ_m
 - Determine nearest scale σ_n
 - Sample Gaussian window, magnitude and angle image at $\sigma_n \rightarrow G_s, M_s, \theta_s$
 - Compute weighted gradient magnitude $M_w = G_s.* M_s$
 - Rotate magnitude angles of θ_s by $\theta_{s,rot} = \theta_s - \theta_m$
 - Build the 16 histograms with 8 bins (each bin covers $360/8 = 45$ degrees)
 - Build descriptor vector \rightarrow rearrange values of Histograms (128 bins) to a vector
 - Normalize the vector
 - **end**
- **end**

Descriptor Computation: Algorithm

- compute the 49×49 Gaussian filter ($\sigma = 8$)
- **for all scales**
 - **for all points in current scale**
 - Initialize arrays x_c, y_c with local image coordinates for a 16×16 neighborhood
 - Rotate local coordinates x_c and y_c using main orientation θ_m
 - Determine nearest scale σ_n
 - Sample Gaussian window, magnitude and angle image at $\sigma_n \rightarrow G_s, M_s, \theta_s$
 - Compute weighted gradient magnitude $M_w = G_s.* M_s$
 - Rotate magnitude angles of θ_s by $\theta_{s,rot} = \theta_s - \theta_m$
 - Build the 16 histograms with 8 bins (each bin covers $360/8 = 45$ degrees)
 - Build descriptor vector \rightarrow rearrange values of Histograms (128 bins) to a vector
 - Normalize the vector
 - **end**
- **end**

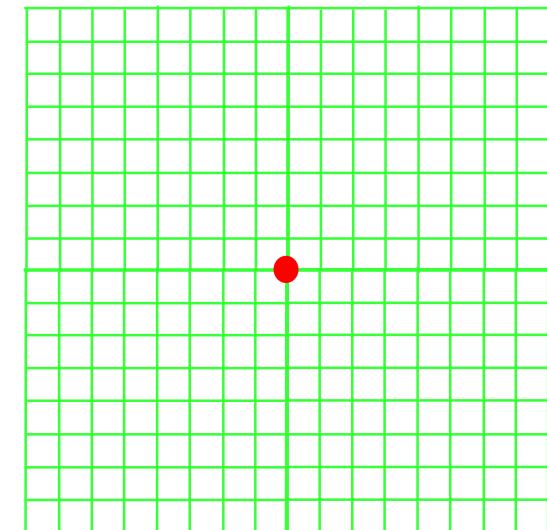
Descriptor Computation 1/5

Local coordinate arrays:

$x_c = repmat([-7.5:7.5], [16,1]);$

$y_c = x_c';$

- Center of descriptor window lays on edge → floating values



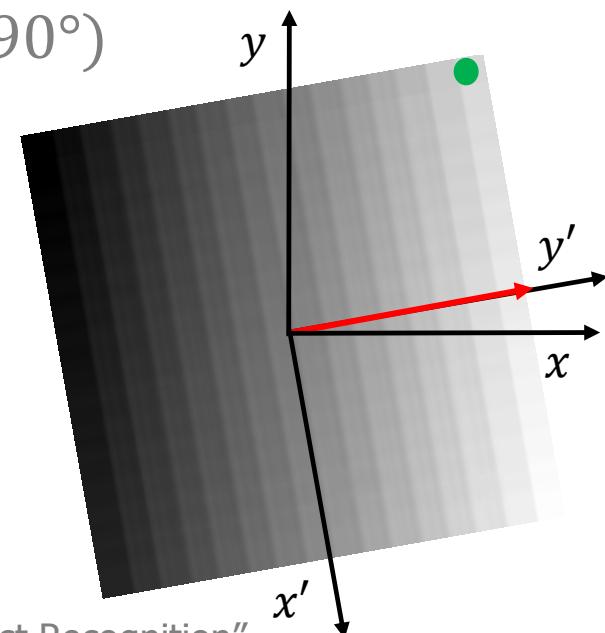
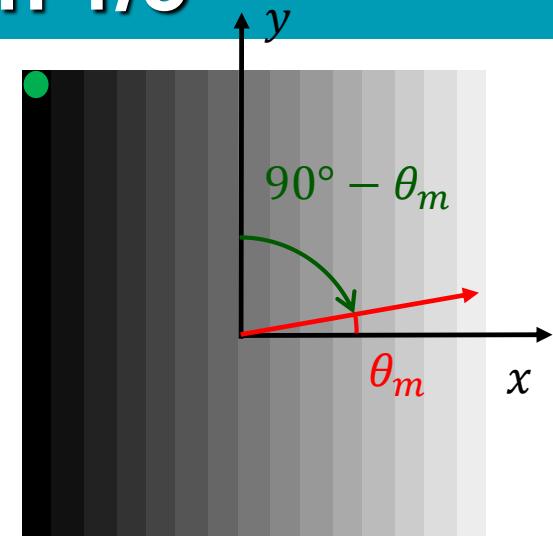
Descriptor Computation 1/5

Rotation Example:

- Rotation of local coordinates x_c when $\theta_m = 10^\circ$ is given

→ Clockwise rotation by $90^\circ - \theta_m$ is desired: $\theta' = -(90^\circ - \theta_m) = (\theta_m - 90^\circ)$

$$\begin{aligned}x'_c &= \sin(\theta') y_c + \cos(\theta') x_c \\y'_c &= \cos(\theta') y_c - \sin(\theta') x_c\end{aligned}$$



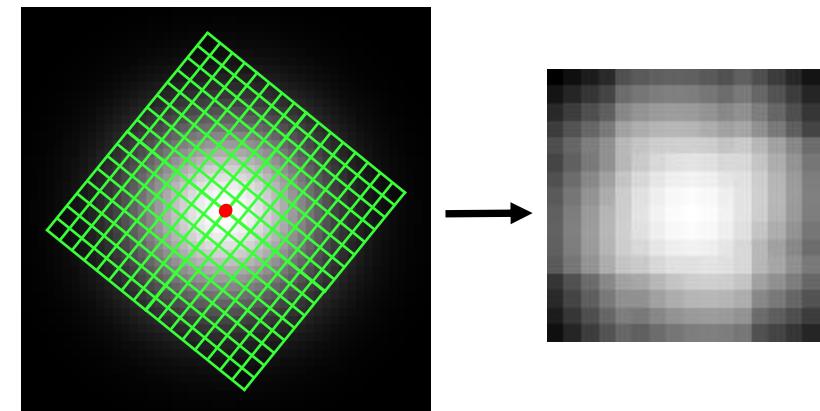
Descriptor Computation: Algorithm

- compute the 49×49 Gaussian filter ($\sigma = 8$)
- **for all scales**
 - **for all points in current scale**
 - Initialize arrays x_c, y_c with local image coordinates for a 16×16 neighborhood
 - Rotate local coordinates x_c and y_c using main orientation θ_m
 - Determine nearest scale σ_n
 - Sample Gaussian window, magnitude and angle image at $\sigma_n \rightarrow G_s, M_s, \theta_s$
 - Compute weighted gradient magnitude $M_w = G_s.* M_s$
 - Rotate magnitude angles of θ_s by $\theta_{s,rot} = \theta_s - \theta_m$
 - Build the 16 histograms with 8 bins (each bin covers $360/8 = 45$ degrees)
 - Build descriptor vector \rightarrow rearrange values of Histograms (128 bins) to a vector
 - Normalize the vector
- **end**
- **end**

Descriptor Computation 2/5

- **Nearest scale** σ_n : Already available
- **Sampling** of Gaussian Window using x_c' and y_c'

- Gaussian Window G (49×49)
 - Translate x_c' and y_c' to center of G
 - here: $\bar{x}_c' = \text{round}(x_c' + 49/2)$;
 - Sample new values G_s
 - Use \bar{x}_c' and \bar{y}_c' as indices
 - Use function *sub2ind*



- **Sampling** of magnitude and angle images using x_c' and y_c'
 - Same concept as for Gaussian window
 - Gradient direction and magnitude images according to σ_n
 - Shift local rotated coordinates \bar{x}_c' and \bar{y}_c' to keypoint position in image → Resampling yields M_s, θ_s

Descriptor Computation: Algorithm

- compute the 49×49 Gaussian filter ($\sigma = 8$)
- for all scales
 - for all points in current scale
 - Initialize arrays x_c, y_c with local image coordinates for a 16×16 neighborhood
 - Rotate local coordinates x_c and y_c using main orientation θ_m
 - Determine nearest scale σ_n
 - Sample Gaussian window, magnitude and angle image at $\sigma_n \rightarrow G_s, M_s, \theta_s$
 - Compute weighted gradient magnitude $M_w = G_s.* M_s$
 - Rotate magnitude angles of θ_s by $\theta_{s,rot} = \theta_s - \theta_m$
 - Build the 16 histograms with 8 bins (each bin covers $360/8 = 45$ degrees)
 - Build descriptor vector \rightarrow rearrange values of Histograms (128 bins) to a vector
 - Normalize the vector
- end
- end

Descriptor Computation 3/5

- **Weighted gradient magnitude:** $M_w = G_s.* M_s$
- **Rotate sampled local gradient directions**
 - Local main orientation θ_m already available

$$\theta_{s,rot} = \theta_s - \theta_m$$

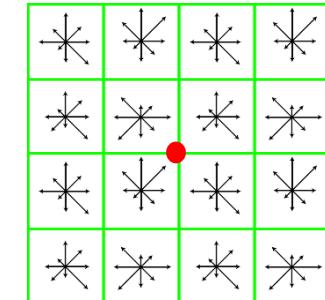
Descriptor Computation: Algorithm

- compute the 49×49 Gaussian filter ($\sigma = 8$)
- **for all scales**
 - **for all points in current scale**
 - Initialize arrays x_c, y_c with local image coordinates for a 16×16 neighborhood
 - Rotate local coordinates x_c and y_c using main orientation θ_m
 - Determine nearest scale σ_n
 - Sample Gaussian window, magnitude and angle image at $\sigma_n \rightarrow G_s, M_s, \theta_s$
 - Compute weighted gradient magnitude $M_w = G_s.* M_s$
 - Rotate magnitude angles of θ_s by $\theta_{s,rot} = \theta_s - \theta_m$
 - **Build the 16 histograms with 8 bins (each bin covers $360/8 = 45$ degrees)**
 - Build descriptor vector \rightarrow rearrange values of Histograms (128 bins) to a vector
 - Normalize the vector
 - **end**
- **end**

Descriptor Computation 4/5

Histogram building

- 16 Histograms of gradient directions
 - One for each 4×4 px cell of $\theta_{s,rot}$
 - 45° binsize → 8 bins



- Use weighted gradient magnitudes instead of adding 1

- Usually

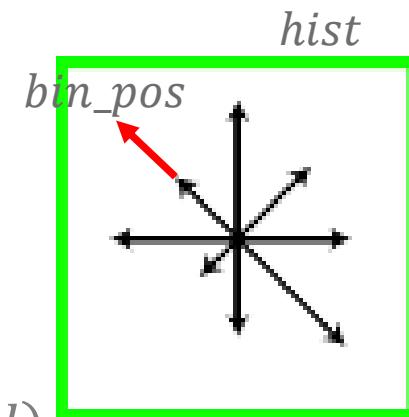
$$hist(bin_pos) = hist(bin_pos) + 1;$$

- has length of 1

- Here

$$hist(bin_pos) = hist(bin_pos) + M_w(row, col)$$

- has length of corresponding weighted gradient magnitude $M_w(row, col)$

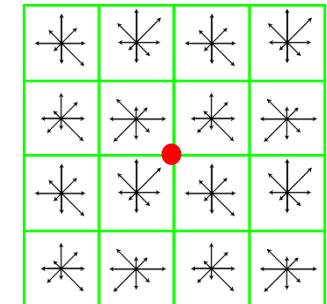


Descriptor Computation: Algorithm

- compute the 49×49 Gaussian filter ($\sigma = 8$)
- **for all scales**
 - **for all points in current scale**
 - Initialize arrays x_c, y_c with local image coordinates for a 16×16 neighborhood
 - Rotate local coordinates x_c and y_c using main orientation θ_m
 - Determine nearest scale σ_n
 - Sample Gaussian window, magnitude and angle image at $\sigma_n \rightarrow G_s, M_s, \theta_s$
 - Compute weighted gradient magnitude $M_w = G_s.* M_s$
 - Rotate magnitude angles of θ_s by $\theta_{s,rot} = \theta_s - \theta_m$
 - Build the 16 histograms with 8 bins (each bin covers $360/8 = 45$ degrees)
 - Build descriptor vector \rightarrow rearrange values of Histograms (128 bins) to a vector
 - Normalize the vector
 - **end**
- **end**

Descriptor Computation 5/5

- Result of histogram building:
 - 16 histograms with 8 bins $\rightarrow 16 \times 8 = 128$ values
- Build a descriptor vector d
 - Rearrange values for each keypoint to one vector d of 128 elements
- Normalize the descriptor to unit length
$$d = \frac{d}{\|d\|} = \frac{d}{norm(d)}$$
 - Reduces effects of illumination change
- Crop all values $d > 0.2 \rightarrow$ set these values to 0.2
 - Reduces non-linear illumination changes
- Normalize d again \rightarrow finished



Matlab Code Structure: Main Function

```
% Main function: TestSIFTFeatures
% Here the whole processing chain is done
% Used functions:
%   - CalcSIFTFeatures to do steps 1-7 for two input images
%   - MatchSIFTFeatures to do step 9
%   - VisualizeMatches to do step 10
function TestSIFTFeatures

    % select and import image data
    [file1, path1, image1, img_size1] = read_image('Select file 1 for processing');
    image1 = mat2gray( image1 );
    [file2, path2, image2, img_size2] = read_image('Select file 2 for processing');
    image2 = mat2gray( image2 );

    % calculate SIFT features
    plot = 1; % 0 --> no plotting of intermediate results; 1 --> plotting activated
    pre_smooth = 1; % 0 --> no prior smoothing; 1 --> prior smoothing (more robustness)
    thresh_contrast = 0.03;%; % the higher, the less points - default: 0.03
    thresh_edges = 6; %the lower, the less points - default: 10

    [Keypoints1, Locations1, image1] = CalcSIFTFeatures(image1, img_size1, plot, pre_smooth, thresh_contrast, thresh_edges);
    [Keypoints2, Locations2, image2] = CalcSIFTFeatures(image2, img_size2, plot, pre_smooth, thresh_contrast, thresh_edges);

    if ~length(Keypoints1(:,1)) | ~length(Keypoints2(:,1))
        return;
    end
    % match points
    thresh_dist = 0.2; % all matches [0,...,sqrt(128)] gt thres_dist will be dismissed
    [matches] = MatchSIFTFeatures(Keypoints1, Keypoints2, thresh_dist);

    % visualize result
    VisualizeMatches(image1, image2, matches, Locations1, Locations2)
end
```

```

function [Keypoints, Locations, Image] = CalcSIFTFeatures(image, img_size, plot, pre_smooth, thresh_contrast, thresh_edges)

%
% Calculate scale space and differences of Gaussians
% caution: pree-smoothing leads to an doubled image size (nn-sampling)
[DoG_pyramid, Sigmas, Image_pyramid, img_size] = CalcDoGPyramid(n_int, image, sigma_start, sigma_stop, plot, pre_smooth);

%
% find initial local maxima and minima in DoG Images
[Maxima, Minima] = FindLocalExtrema(DoG_pyramid, Sigmas);

%
% refine the results

% at first just do shifting of extrema, if position tells us to do so
shifting = 1; % therefore, just set shifting to 1
% this avoids the refinement and just shifts yx positions if necessary
Maxima = LocalizeContrastEdges(Maxima, DoG_pyramid, Sigmas, thresh_contrast, thresh_edges, shifting);
Minima = LocalizeContrastEdges(Minima, DoG_pyramid, Sigmas, thresh_contrast, thresh_edges, shifting);

% now, we can do refinement with new positions (contrast and edge-iness)
shifting = 0; % no shifting, just filtering out due to contrasttt and edge-iness
Maxima = LocalizeContrastEdges(Maxima, DoG_pyramid, Sigmas, thresh_contrast, thresh_edges, shifting);
Minima = LocalizeContrastEdges(Minima, DoG_pyramid, Sigmas, thresh_contrast, thresh_edges, shifting);

%
% orientation assignment: compute orientations for each scale, where
% extrema occur
[Magnitude_Pyramid, Angle_Pyramid] = ComputeOrientations(Image_pyramid, Maxima, Minima);

%
% assign orientations
[Maxima] = AssignOrientations(Magnitude_Pyramid, Angle_Pyramid, Maxima, Sigmas);
[Minima] = AssignOrientations(Magnitude_Pyramid, Angle_Pyramid, Minima, Sigmas);

%
% plot results
if plot
    title = 'Refined maxima (red) and Minima (green) with orientation and magnitude.';
    PlotExtrema(Maxima, Minima, Image_pyramid{1}, title);
end
%

% We are not longer interested whether the points are max or mins
% --> concatenation to one array each scale
for s = 1:length(Maxima)
    Extrema{s} = [Maxima{s}; Minima{s}];
end

%
% descriptor computation
[Keypoints, Locations] = ComputeDescriptor(Magnitude_Pyramid, Angle_Pyramid, Extrema, Sigmas, img_size);

```

Task: Implement Function *ComputeDescriptor*

```
% descriptor computation  
[Keypoints, Locations] = ComputeDescriptor(Magnitude_Pyramid, Angle_Pyramid, Extrema, Sigmas, img_size);
```

Provided:

```
%-----  
% Computation of the final descriptor  
% Inputs:  
% - Magnitude_Pyramid: Cell with length(Magnitude_Pyramid) = length(Maxima)  
%   --> in each level for each pixel the local gradient magnitudes are stored  
% - Angle_Pyramid: Cell with length(Angle_Pyramid) = length(Maxima)  
%   --> in each level for each pixel the local gradient orientations are stored  
% - Maxima: Cell with maxima in each scale step-->array of size (n_points, 5)  
%           [y,x,sigma,orientation,magnitude]  
% - Minima: see maxima  
% - Sigmas: vector of size (n_scales) with corresponding values sigma for  
%           each scale level of DoG-pyramid  
% - img_size: size of the image [y,x]  
% Outputs:  
% - Keypoints: Array with keypoint information for the image  
%               Dimensions: [n_keypoints, 128]  
% - Locations: Coordinates of keypoints for the image; in  
%               each row of that array (size [n_points, 5]) we have:  
%               [y,x,sigma,orientation,magnitude]  
% The Keypoits- and the Locations-array will have the same number of rows  
% and a row in Keypoits(i) corresponds to a row Locations(i) !  
% After computation there will be not longer a separation of minima and  
% maxima!  
function [Keypoints, Locations] = ComputeDescriptor(Magnitude_Pyramid, Angle_Pyramid, Extrema, Sigmas, img_size)  
end
```

Implement descriptor computation here!

ComputeDescriptor: Inputs

```
% descriptor computation  
[Keypoints, Locations] = ComputeDescriptor(Magnitude_Pyramid, Angle_Pyramid, Extrema, Sigmas, img_size);
```

- ***img_size***: 2-element vector with image size (row,col)
- ***Sigmas***: $1 \times n_s$ vector with σ for each scale of the pyramid
 - n_s = number of scale space steps
- ***Magnitude_Pyramid***: $1 \times n_s$ cell with gradient magnitudes
 - *Magnitude_Pyramid*{1} contains an array of *img_size* with gradient magnitudes in scale *Sigmas*(1)
- ***Angle_Pyramid***: $1 \times n_s$ cell with gradient orientations
- ***Extrema***: $1 \times n_s$ cell with keypoints for each scale

ComputeDescriptor: Inputs

```
% descriptor computation  
[Keypoints, Locations] = ComputeDescriptor(Magnitude_Pyramid, Angle_Pyramid, Extrema, Sigmas, img_size);
```

Extrema: $1 \times n_s$ cell with keypoints for each scale

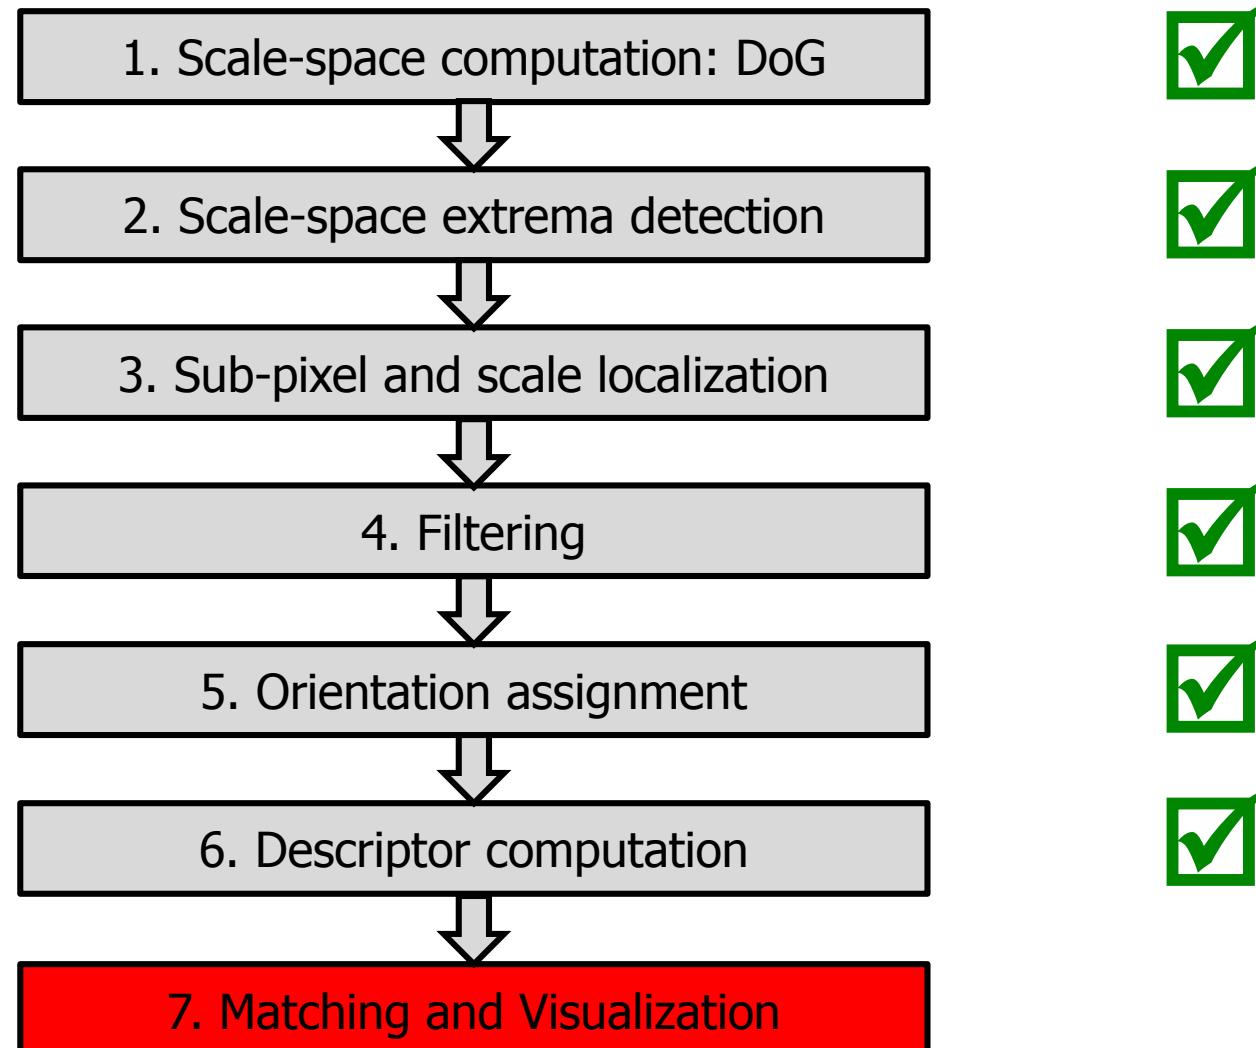
- Get extrema for scale σ : $ex = Extrema\{\sigma\}$;
 - ex is an array of size $n \times 5$, where n is the number of extrema
 - If there are no extrema in scale σ , ex will be an empty array
 - Otherwise, each row r of ex obtains the following information:
 - $ex(r, :) = [y, x, \sigma, \theta_m, m_m]$
 - y, x : subpixel position of keypoint ($y \rightarrow$ rows, $x \rightarrow$ columns)
 - σ : estimated scale of keypoint r
 - θ_m : Main gradient direction
 - m_m : Main gradient magnitude (not needed in computations, just for visualization)

ComputeDescriptor: Outputs (Your Task)

```
# descriptor computation  
[Keypoints, Locations] = ComputeDescriptor(Magnitude_Pyramid, Angle_Pyramid, Extrema, Sigmas, img_size);
```

- **Keypoints:** array of size $n_p \times 128$ with the descriptor vectors
 - n_p is the overall number of keypoints (scale-independent!)
 - Each row represents the descriptor vector of one keypoint
- **Locations:** array of size $n_p \times 5$ with corresponding information for each keypoint
 - Store $ex(r, :) = [y, x, \sigma, \theta_m, m_m]$ here (also scale-independent)
 - Rearrange input values of *Extrema*
- Order of Keypoints and Locations has to be the same!

Algorithm Outline



Matching and Visualization

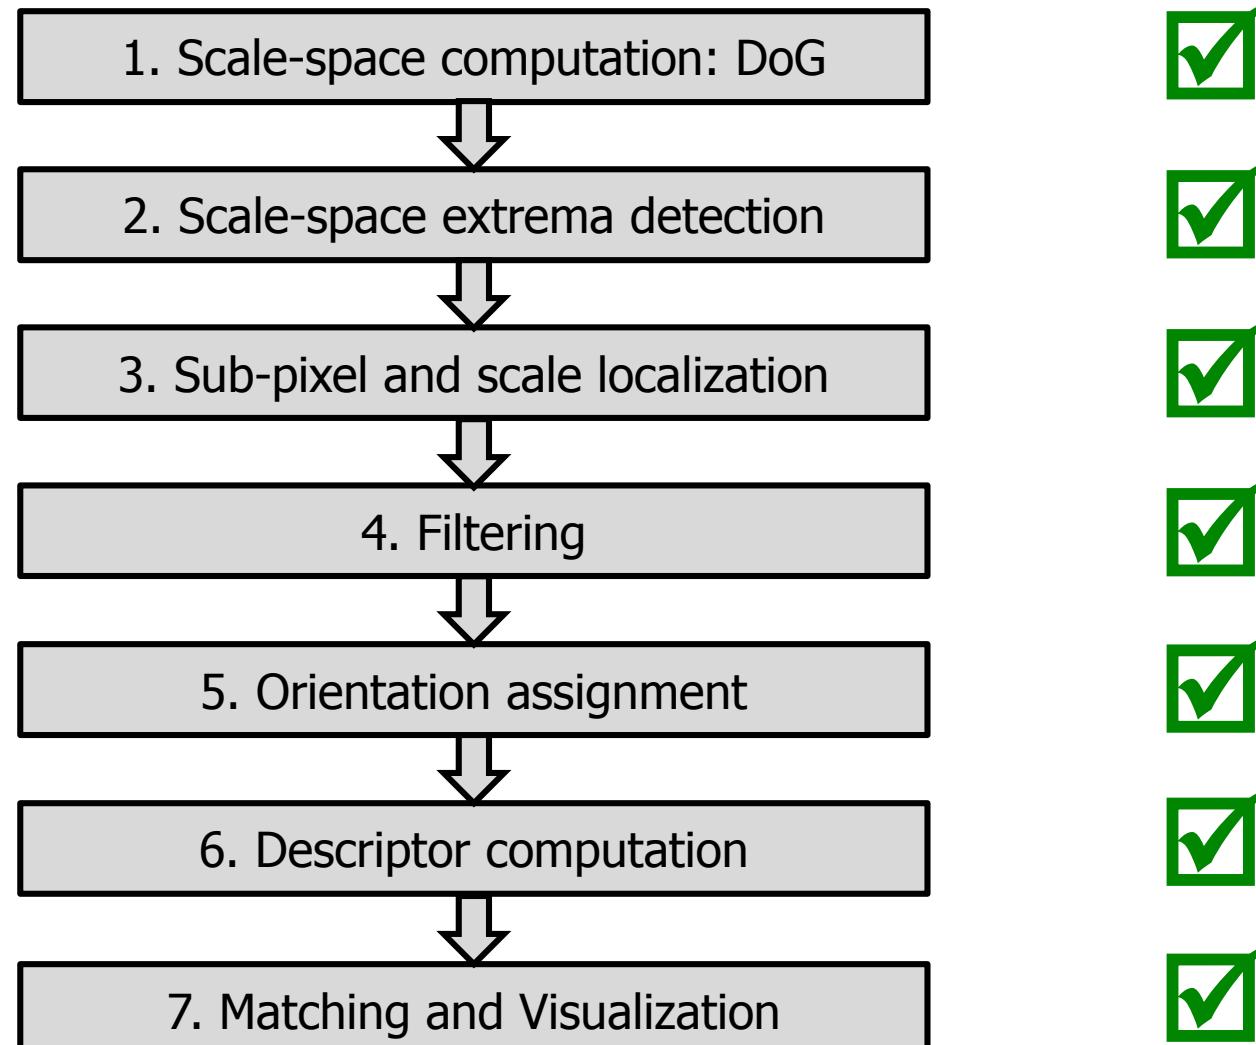
Comparison of all $i = 1, \dots, n_p$ descriptors from image 1 and 2
($descr_i^1$ and $descr_i^2$)

- Euclidean norm for distance measurement

$$\begin{aligned} d_i &= \sqrt{\sum_{j=0, \dots, 128} (descr_{i,j}^1 - descr_{i,j}^2)^2} \\ &= \text{norm}(descr_{i,j}^1 - descr_{i,j}^2) \end{aligned}$$

- Eliminate double matches

Algorithm Outline



Example

- Input Images

Input1.jpg

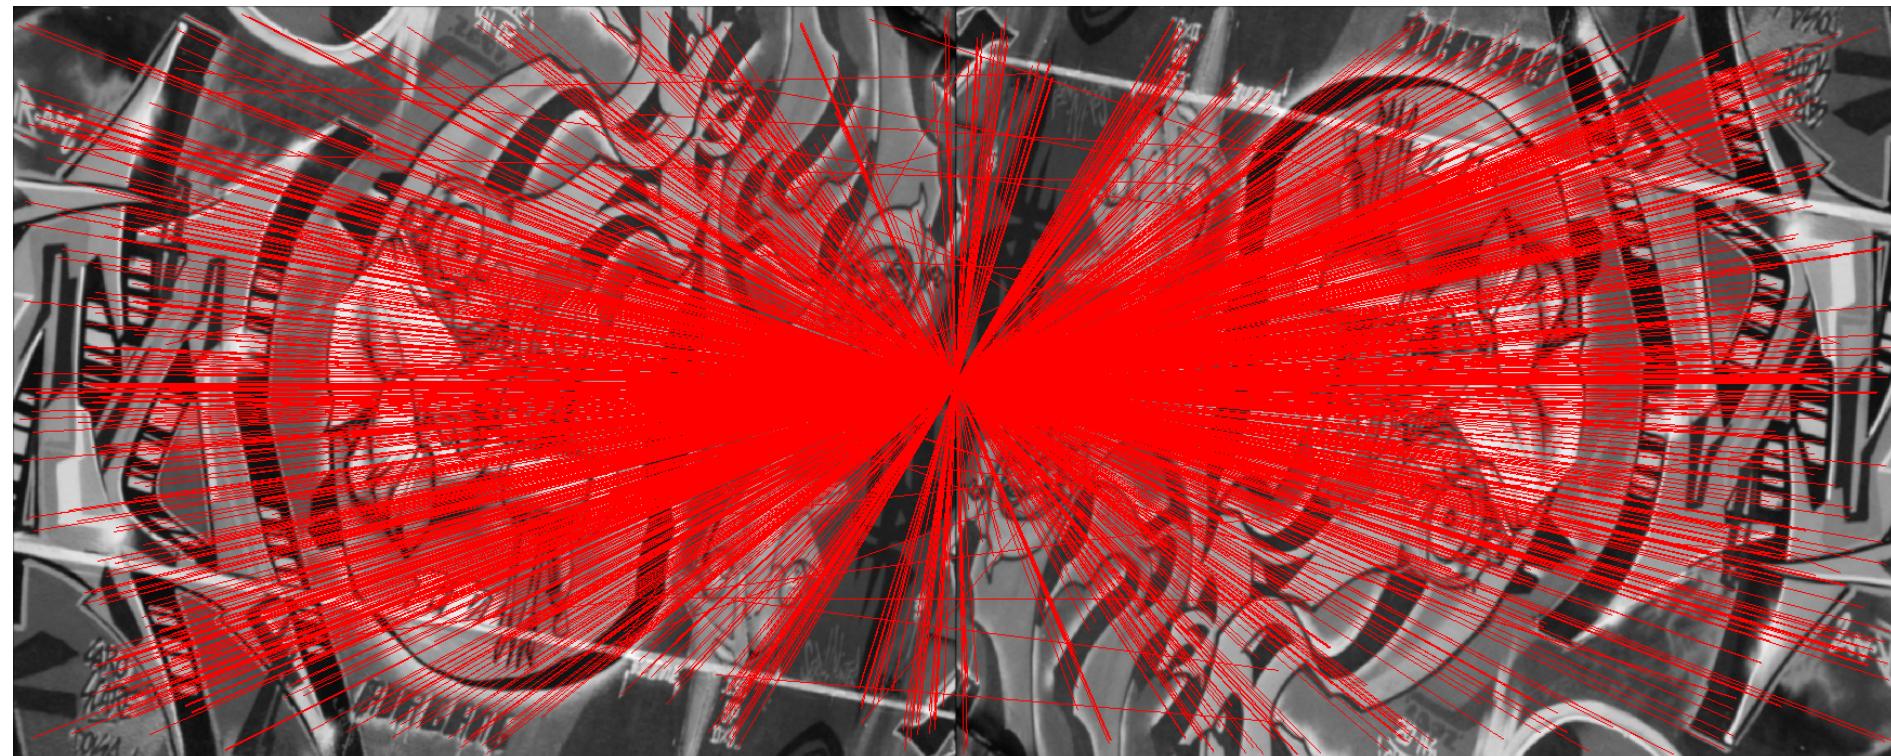


Input2.jpg



Example

- Desired Result



- Tipp: Use subimages for testing (runtime)

Provided Data (See CV-Page)

- Matlab source: *TestSIFTFeatures.m*
- Octave source: *TestSIFTFeatures_octave.m*
- Full-size images: *Input1.jpg*, *Input2.jpg*



- Small-size images: *Input1_s.jpg*, *Input2_s.jpg*





Thank you!